i have an idea for an app/site. help me flush it out and research if its already being done

May 24

Claude responded: What's the idea?

What's the idea? Tell me as much or as little as you have — even a rough concept is enough to start with.

You said: it would be a crowdsourcing image and video app/site that would allow people to upload media from specific events they attend.

it would be a crowdsourcing image and video app/site that would allow people to upload media from specific events they attend. so, everyone who attended a concert could upload their pics and videos from that concert for everyone else to view and enjoy. the way i'm thinking about it would be each event/artist/team etc would have their own dedicated page (cardi b page, kentucky derby page, philadelphia phillies page, indy 500 page, etc etc) and then within each page there could be individual 'sub' events (say, the various concerts that cardi b has, or the 2025 kentucky derby and then the 2026 kentucky derby, etc) so people would attach/upload their media that they took while at those events into the correct show/time. the idea would be to allow everyone else (who either did or did not attend the event) to view all the cool media that was taken by event attendees.

May 24

Claude responded: Really solid concept.

Really solid concept. Let me research the competitive landscape first, then we can flush out the idea properly.

Searched the web

Searched the web

Good research. Here's a full breakdown:

The Concept, Fleshed Out

Core idea: A public, permanent, searchable archive of fan-captured media — organized by entity (artist/team/event brand) → specific occurrence (date/show/game). Think of it as **the fan's version of Getty Images** — but free, crowdsourced, and built around the experience of being there.

What Already Exists (Competitive Landscape)

The space has been tried before, which is actually validating — but none of the attempts nailed what you're describing:

CrowdAlbum (2013–2016) — This was the closest match to your idea. It aggregated photos and videos from performances, helped fans live vicariously, and gave artists a visual history of their tours. Spotify acquired it in 2016, then quietly killed it. The gap it left has never been properly filled. Android PoliceAndroid Police

CrowdFlik — Concertgoers checked in to an event, uploaded video clips, and the app organized shared videos based on location and time — enabling users to create a composite video ("flik") from crowd-sourced footage. More of a video mashup tool than a media archive. BizBash

SureShot — An AI platform that collects crowd-generated videos from concerts, festivals, and sports events. But it's **B2B** — event organizers pay to collect UGC for their own promotional use. Not a public fan-facing archive. SureShot

Happaning — Transforms fan videos into immersive experiences with AI-powered insight, focused on helping brands, festivals, and artists capture and own content their audience is already creating. Again — B2B, not fan-community-forward. Happaning

GuestCam / Fotify / Kululu — Event photo sharing apps focused on weddings, corporate events, and private gatherings — not public live events or sports. Fotify

The Gap You're Identifying

Everything out there is either:

- **Dead** (CrowdAlbum)
- **B2B / organizer-facing** (SureShot, Happaning) — fans contribute but don't own the experience
- **Private event focused** (GuestCam, Fotify) — not concerts/sports
- **Social media** (Instagram, TikTok) — content is scattered, unsearchable by event, and

disappears into feeds

Nobody owns the public fan archive — the permanent, browsable, community-owned record of "I was at the 2025 Kentucky Derby and here's what it looked like from the stands."

Key Differentiators to Build Around

1. **Entity → Event hierarchy** (exactly as you described) — Artist/Team/Venue page → individual show/game sub-pages. This is the core architecture that nobody has done cleanly for the public.
2. **Fan-first, not brand-first** — Unlike SureShot/Happaning, the content belongs to the community, not the promoter.
3. **Permanent archive** — Not a feed that scrolls away. The 2019 World Series should still be browsable in 2035.
4. **Cross-category** — Concerts, sports, motorsports, rodeos, air shows, parades. Anything with an audience and a specific date.
5. **Discovery angle** — "I'm going to Coachella next month — let me see what it looked like last year from the crowd's perspective."

Potential Monetization

- **Verified/official partnerships** — Teams and artists pay to feature or curate their pages
- **Licensing** — Fans can optionally license their media; platform takes a cut
- **Premium uploader tiers** — Higher upload limits, video quality, etc.
- **Advertising** — Contextual ads on high-traffic event pages
- **Data/analytics** — Sell aggregated fan attendance/engagement insights to venues and promoters (CrowdAlbum was heading this direction before Spotify grabbed them)

Biggest Challenges to Think Through

- **Rights/IP** — Concert venues and sports leagues have strict photography policies. You'd need clear ToS about what fans can legally upload and what the platform can do with it.
- **Content moderation** — At scale, UGC is messy.
- **Cold start problem** — Pages with no content attract no uploaders. You'd need a seeding strategy (pull from Instagram/TikTok hashtags? Partner with a few artists early?)
- **Video storage costs** — Video at scale is expensive. You'd need a clear strategy (compression, tiered storage, limits on video length).

Bottom Line

The idea is **legitimately unoccupied territory** at the consumer/community level. The closest thing that existed (CrowdAlbum) was acquired and killed nearly a decade ago, and nothing has replaced it for fans. The B2B players (SureShot, Happaning) have validated demand from the supply side but left the demand side (fans who want to browse and relive events) completely unserved.

Want to go deeper on any piece of this — architecture, go-to-market, monetization, or tech stack?

You said: yes.

yes. go deeper on all of it

May 24

Claude responded: Let me do a bit more targeted research first before we build this out fully.

Let me do a bit more targeted research first before we build this out fully.

Searched the web

Searched the web

Good. I have enough. Let me build this out comprehensively.

Now let me build this out. This will be a comprehensive deep-dive across all dimensions.

V

visualize

V

visualize show_widget

Here's the full deep-dive across every dimension. Click any node in the diagram above for a quick drill-down.

1. Product architecture

The three-tier hierarchy

The core data model is simple but the execution matters:

Tier 1 — Entity page (e.g. "Cardi B", "Philadelphia Phillies", "Kentucky Derby") This is a permanent, evergreen page. It has a bio/about section, a stats panel (total uploads, events covered, top contributors), and a chronological list of all events. Entities can be claimed/verified by the official rights holder down the road — but they exist with or without official involvement. Think Wikipedia-style: fans create it, officials can enhance it.

Tier 2 — Event page (e.g. "Cardi B — Chase Center, San Francisco — Oct 14 2025") Each event gets a unique, permanent URL. Key fields: artist/entity, venue, city, date, optional tour name. This page is the unit of organization — every piece of media attaches here. Events can be pre-created (pulled from Ticketmaster/setlist.fm APIs before the event happens) so fans arrive at a page that already exists.

Tier 3 — Media grid The browsable gallery of everything uploaded to that event. Supports photos and video. Sortable by: newest, most liked, "staff picks." Filterable by: photos only, videos only, section of venue (if you build that feature). Each media item has uploader credit, timestamp, optional caption.

The upload flow

This is the make-or-break UX moment. It has to be nearly frictionless:

1. User scans a QR code at the venue (printed on tickets, venue signage, or you push a notification if they've used the app before) — OR they find the event page on their own after the fact
2. One-tap to select from camera roll
3. Tag to the correct event (pre-populated if they're at the venue via geolocation, or searched)
4. Optional: add caption, tag their seat section
5. Upload queues in background — they can leave the app

No account required for upload — allow anonymous guest uploads initially to maximize participation. Accounts unlock: seeing all your past uploads, getting notified when your uploads get liked, building a profile, higher upload limits.

2. Event database strategy

You don't build this from scratch. You ingest from existing sources:

Ticketmaster already integrates setlist.fm content directly into Artist Detail Pages via API — meaning there's a rich, publicly accessible ecosystem of event data you can tap. Specifically: [Ticketmaster Business]

- **Ticketmaster Discovery API** — free tier, covers upcoming events globally with venue,

date, artist, tour name

- **Setlist.fm API** — returns structured JSON with artist, venue, city, date, tour name, and setlist for past events — invaluable for backfilling historical events and giving pages immediate content Setlist
- **Sports-specific**: ESPN API, MLB/NFL/NBA official APIs, or third-party aggregators for game schedules and results

The workflow: ingest event data nightly via cron job → auto-create event pages → enrich with setlist data after the show → allow fan uploads from the moment the event is announced. This means by the time someone searches for "Beyoncé Miami 2025," the page already exists and is waiting for their photos.

3. Tech stack

Frontend

React or Next.js — you need SSR for SEO. These event pages need to rank on Google ("Cardi B concert photos 2025") and server-side rendering is non-negotiable for that. Next.js with App Router is the current standard choice.

Mobile app: React Native shares a ton of logic with the web codebase. The mobile app is important for the upload experience — you want a native camera roll integration, not a web upload picker.

Media storage

This is the single most consequential technical decision. Cloudflare R2 at $0.015/GB stored with free egress is 60-80% cheaper than S3+CloudFront at scale, and the delivery cost of S3+CloudFront is roughly 18x the storage cost — so egress is where you'd get killed if you chose wrong. R2 is the right call for a media-heavy platform. LeanOps TechnologiesLeanOps Technologies

For video specifically: Mux charges ~$0.007/min delivered but includes encoding, adaptive bitrate streaming, and analytics — worthwhile for a platform where video quality matters. You get HLS adaptive streaming, thumbnail generation, and a player that works everywhere. LeanOps Technologies

Search

Algolia or Typesense for event/entity search. You need instant, typo-tolerant search ("Beyoncé" vs "Beyonce", "Phillies" vs "Phils"). Algolia has the better product; Typesense is cheaper and self-hostable. Either integrates well with a Next.js stack.

Database

PostgreSQL for structured data (entities, events, users, metadata). The schema is fundamentally:

- `entities` (id, name, type, slug, verified)
- `events` (id, entity_id, venue, city, date, tour_name, external_id)
- `media` (id, event_id, user_id, url, type, likes, created_at)
- `users` (id, username, upload_count, etc.)

4. The legal landscape

This is real but manageable. Here's the actual picture:

Fan photography rights: In the U.S., public photography and filming are largely protected under First Amendment rights, and venues reinforce this with ticket disclaimers and signage, legally framing attendance as consent. The practical reality is that fans film everything at concerts and sports events and nobody is prosecuted for it. Fairpicture

Copyright ownership: The photographer (not the artist, venue, or promoter) owns the copyright in the photos they shoot. So fans uploading their own photos own those photos. Your platform is a host, not a rights holder. Cassels

Your platform's liability: The DMCA safe harbor protects you as a platform hosting user content, as long as you have a proper DMCA takedown process, don't have actual knowledge of infringement, and respond promptly to valid takedown requests. This is the same protection that covers YouTube, Instagram, and every other UGC platform. You're not in unusual legal territory here.

What you need in your ToS:

- Users warrant they own or have rights to what they upload
- Users grant the platform a license to display, distribute, and create thumbnails/previews
- Clear DMCA contact and process
- No commercial use of others' uploads without explicit license

The one gray area: Some venues and artists explicitly prohibit video recording in their ticketing terms. Users violating those terms is a contractual issue between them and the venue — it doesn't make your platform liable for hosting the content, assuming you have a good DMCA process. That said, a clear community guideline that users should follow venue rules, and a mechanism for venues to request takedowns of specific recordings, will help you maintain good relationships with the industry.

5. Go-to-market strategy

Cold start solution

The empty-page problem is real — nobody uploads to a page with nothing on it. The fix is a two-pronged seeding strategy:

Historical backfill: Use setlist.fm + Instagram/Twitter scraping (public hashtags like

`#beyonceconcert`, `#2025WorldSeries`) to seed pages with existing fan content that's already public. You're not hosting it — you're embedding/linking. This creates the appearance of a live community on day one. Over time, original uploads replace embeds.

Partner early with 3-5 mid-tier artists or teams who have passionate fanbases and will promote the platform from their own channels. Not stadium-level acts (they have lawyers and IP concerns). Think: touring bands with 100k-500k followers, minor league sports teams, regional racing circuits. Their fans are hungry for community and not overprotected by label/league IP enforcers.

Distribution

- **SEO is your primary acquisition channel.** Every event page is a long-tail keyword that nobody is ranking for. "Taylor Swift Eras Tour Philadelphia night 2 fan photos" has search volume and zero strong results. You want to be the answer. This means clean URLs, proper metadata, structured data markup for events.
- **Social seeding**: Every photo/video uploaded should have a shareable card with your branding. When fans share their uploads to Instagram or TikTok, that's a distribution touchpoint.
- **Ticket integrations**: Eventually you want a Ticketmaster or SeatGeek API relationship where you can surface "see fan photos from this event" on the ticketing page. That's a moat.

6. Monetization (in order of execution)

Phase 1 — Free, build the network No monetization for the first 12-18 months. The goal is uploads, not revenue. Every dollar gate you add at this stage costs you 10x in growth.

Phase 2 — Premium uploader tier ($5-10/month) Higher upload limits (free: 20 photos + 2 videos per event; premium: unlimited), HD video uploads, ability to mark uploads as "licensed" for others to request permission to use, early access to features. This is your first revenue line and it targets the power users who are already uploading a lot.

Phase 3 — Official entity partnerships Teams, artists, and venues pay for a "verified" page experience: curated official content mixed in with fan uploads, custom branding on the page, ability to feature certain uploads, access to engagement analytics. Pricing: $500-2,000/month per entity, or a rev-share on licensed content. This is B2B SaaS on top of a consumer platform — a proven playbook.

Phase 4 — Content licensing marketplace Fans can list their uploads as available for licensing. Media outlets, brands, documentarians can request a license. Platform takes 20-30% of the transaction. This is where a single viral video from a historic event becomes real revenue for the uploader — and a powerful word-of-mouth story for the platform.

Phase 5 — Contextual advertising High-intent, low-friction. Someone browsing photos from a Phillies game is a perfect audience for stadium food delivery, sports merchandise, or the team's streaming package. CPMs on sports/entertainment UGC pages will be strong. Keep ads sparse

and non-interruptive — the community will revolt at banner-heavy pages.

7. Competitive moat over time

The business gets stronger as it scales in ways competitors can't easily replicate:

1. **Data network effect** — every upload makes every page better. A competitor starting from zero can't replicate 3 years of fan uploads across 50,000 events.
2. **SEO compounding** — each event page accumulates backlinks, organic traffic, and content over time. A 2019 World Series page keeps getting traffic forever.
3. **Uploader identity** — fans who've built a profile, accumulated likes, and established a reputation on your platform have real switching costs.
4. **API relationships** — once you have partnerships with Ticketmaster, setlist.fm, or sports leagues, those are hard to displace.
5. **Archive positioning** — if you market yourself as "the fan's visual record of live events," you have a cultural mission that goes beyond just another photo app. That framing matters for press, for partnership conversations, and for user loyalty.

8. Name / brand direction

A few directions worth considering:

- **Archive-forward**: Fanarchive, Crowdview, Showvault — emphasizes permanence and the historical value
- **Fan-forward**: Fanframe, Crowdshot, Showside — emphasizes community and perspective
- **Experience-forward**: Wewere.there, Seatview, Livefeed — emphasizes "I was there"

The strongest branding territory is something that communicates **you were there and here's the proof** — both for the uploader and the viewer who wishes they had been.

Want me to go deeper on any specific piece — the database schema, the mobile upload UX, the SEO strategy, building the MVP, or the partnership pitch deck?

You said: Tell me more about Entity pages

Tell me more about Entity pages

May 24

Claude responded: Great drill-down.

Great drill-down. Let me think through Entity pages fully — this is actually the most important page type on the platform, because it's the first thing a fan, an artist's team, or a potential partner lands on.

Entity pages — full breakdown

What is an "entity"?

An entity is any **recurring subject** that generates live events over time. The platform needs to support multiple types cleanly:

Type	Examples
Music artist	Cardi B, Phish, Metallica, a local band
Sports team	Philadelphia Phillies, Duke Blue Devils, Miami Heat
Recurring event brand	Kentucky Derby, Indy 500, Coachella, Super Bowl
Venue	Madison Square Garden, Red Rocks Amphitheatre
Tour	Taylor Swift Eras Tour, Beyoncé Renaissance Tour

Venue and tour pages are secondary but important — a fan might search "Red Rocks 2025" and want to see everything shot there across all artists that year.

The anatomy of an entity page

V

visualize

V

visualize show_widget

The mockup above is interactive — click the tabs (Events, Top media, Contributors, About) to see each section. Here's the thinking behind every piece of it.

Hero section

The header establishes credibility immediately. A few things worth calling out:

The stats bar is load-bearing. "84 events archived, 12,480 photos" tells a new visitor this place is alive and worth their time before they scroll a single pixel. These numbers need to update in real time as uploads come in. For sports teams especially, archive depth is a flex — "We have every Cubs home game since 2019" is a community achievement worth displaying.

Verification is a spectrum, not binary. You'll have three states:

- Unverified — community-created, anyone can suggest edits
- Claimed — the artist/team/org has authenticated and can edit the bio and hero image, but fan uploads remain community-controlled
- Official partner — paying partner, gets enhanced branding, analytics access, ability to feature picks

The "claimed but not partner" tier is important — give official parties control over their page's presentation for free. That goodwill leads to paid partnerships later.

Follow is a deceptively powerful feature. A fan following Cardi B gets notified when a new event page is created (upcoming show announced), when uploads start coming in after a show, and when their own uploads hit milestones. This is your retention and re-engagement loop. Follows also give you a dataset — you know exactly which entities have the most passionate audiences, which is valuable for partnership conversations.

Events tab — the core value

The events list is the heart of the page. A few design decisions here that matter:

Upcoming events show before past. When someone lands on this page before a show they're attending, you want them to immediately see the upcoming event, tap into it, and understand that this is where they should upload after. Pre-show anticipation is a powerful behavior trigger.

Thumbnail strips on past events. Four preview thumbnails on each event row are doing heavy lifting — they communicate "this is a visual archive" without requiring a click. A row with zero uploads gets a muted "no uploads yet" state, which is an implicit prompt for people who were there.

Upload count as social proof. "345 uploads" next to an event row makes the next fan more likely to add theirs. Nobody wants to be the only person who uploaded from that show.

Filters you'll eventually want on this tab: by year, by tour, by city, by "has media" vs. all. Power users browsing a discography will want to jump to a specific era.

Top Media tab

This is the discovery and delight tab — the best stuff surfaced across all events, all time. Algorithmically ranked by likes + view time + recency weighting. This tab is what casual

visitors browse. It's also what press or artists find when they Google the entity page and land here — it's your showcase.

An important decision: **what makes something a "top" upload?** Pure like count is gameable and favors old content. A better signal is: likes per view within the first 48 hours of upload, combined with a time-decay factor. This rewards genuinely great shots from recent shows while letting all-time classics stay visible with a separate "all-time" toggle.

Contributors tab

This is underestimated. The people uploading the most content to a Cardi B page are her most engaged fans — exactly the people her management team would want to identify. This tab:

- Rewards your best uploaders with public recognition
- Creates a status ladder (top contributor badges) that motivates more uploads
- Is a genuine product differentiator for partnership conversations — "we can tell you who your top 50 superfans are in Miami"

The contributor leaderboard should be per-entity and per-event. Being the top contributor at a specific show is a point of pride worth displaying.

About / verification flow

The About tab does two things: shows data provenance (where event data comes from) and handles the claim/verification flow.

The verification request should be lightweight — an email form, a social account verification, maybe a blue-check-style DM from the official account. You're not verifying identity for legal purposes, just establishing that someone from the team has authenticated. Keep the bar low initially to maximize uptake among smaller artists and teams.

Entity types behave differently

The same page template works across entity types, but the content varies:

Music artist pages are tour-centric. The event list groups naturally by tour. Stats that matter: total shows archived, rarest show (small venues from early in their career are gold to superfans).

Sports team pages are season-centric. Events group by season and then by game type (regular season, playoffs, championships). A championship run's event list is a trophy case.

Recurring event brands (Kentucky Derby, Indy 500, Coachella) are year-centric. Each year is essentially a sub-entity. The 2019 Kentucky Derby and the 2025 Kentucky Derby are different events but share the same entity page. Browsing a decade of Derby photos is a compelling experience.

Venue pages are a bonus entity type — "Red Rocks Amphitheatre" aggregates across all artists who've played there. These become SEO workhorses ("Red Rocks 2025 concert photos") and are useful for fans who care more about the place than any specific act.

What you don't show

Just as important as what's on the page:

- No comments on entity pages (too much moderation burden at this level; comments live on individual media items or event pages)
- No edit history visible to regular users (Wikipedia-style, but keep it in the admin layer)
- No personal information about artists beyond what's publicly known
- No setlists by default (you can integrate setlist.fm data, but it's secondary to the media — don't clutter the event rows with song lists unless clicked into)

Want to go into Event pages next, or dive into the upload flow UX?

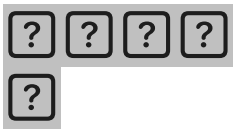

You said: my thinking about this is basically i want the audience (the people who did NOT attend the event) to thunk of this site as a way to essentially 'attend' the ev…

my thinking about this is basically i want the audience (the people who did NOT attend the event) to thunk of this site as a way to essentially 'attend' the event through the uploaded media - maybe they couldn't attend because ticket prices were too high, or they live too far away, etc etc. i'm thinking of it almost as a Youtube-like experience where you would go visit the site not just to get the vibe of the event, but also to essentially virtually attend the event, show, game, etc. thoughts?

May 24

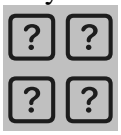

Claude responded: This reframe is significant — and it's actually a stronger product vision than what we had before.

This reframe is significant — and it's actually a stronger product vision than what we had before. Let me think through what it changes.

The original framing was essentially **"fan photo archive."** What you're describing is **"virtual attendance."** Those are meaningfully different products with different design priorities, different content hierarchies, and a much bigger addressable audience.

Here's why this matters:

The audience math flips completely

In the archive framing, your primary user is someone who *was* at the event — they upload their photos, browse others', relive the night.

In the virtual attendance framing, your primary user is someone who **wasn't there** — and that's a vastly larger group. For a Taylor Swift Eras Tour show at a 70,000-seat stadium:

- ~70,000 people attended
- Millions wanted to and couldn't

The people who *couldn't* go are your biggest audience, and they have the strongest emotional motivation to consume content. Ticket prices for major events have become genuinely unaffordable for most people. A Beyoncé floor ticket in 2025 was $800+. The Super Bowl is $10,000+. The Kentucky Derby is $2,000+ for a decent seat. You're building the experience for everyone who got priced out — that's a massive, underserved, emotionally motivated audience.

What this changes about the product

This reframe has real implications for every layer of the platform:

1. Content hierarchy inverts

In an archive, photos are the star. For virtual attendance, **video dominates.** Photos give you a vibe; video puts you there. The YouTube comparison you drew is exactly right — the experience of watching 30 fan-shot videos from different angles of a concert, stitched together mentally into a full show, is genuinely immersive in a way that a photo grid isn't.

This means:

- Video should be the default view, not a filter
- The media grid should lead with video, photos secondary
- Upload limits for free users should be more generous on video than photos
- The playback experience needs to be smooth and full-screen capable — not an afterthought

2. The event page becomes a "show experience" page

V

visualize

V

visualize show_widget

This is what the virtual attendance framing unlocks on an event page. A few things in there worth unpacking:

The three views: Watch, Browse, Map

The toggle in the hero is doing a lot of work:

Watch mode is the default for non-attendees. It's a curated, sequential experience — fan highlights up top, then setlist-linked clips, then a "best of each angle" section. The goal is to make you feel like you watched the show, not just scrolled a photo album. Think of it as an auto-assembled documentary from 94 different phones.

Browse mode is for people who want to explore freely — the full grid, filterable by photo/video, sortable by likes or newest. This is more the archive experience, and it's still valuable, but it's not the lead.

Map mode is a venue seating map where you can tap a section and see all uploads from that vantage point. This is genuinely powerful for two reasons: a non-attendee gets to experience the show from *every* seat in the house, and a future attendee can preview what their seats will look like before they buy. That second use case is a real product hook — "should I buy floor or section 112 tickets? Let me see what section 112 looks like at this venue for this type of show."

Setlist-to-clip linking

This is the feature that makes the virtual attendance experience actually coherent. Right now if you want to watch fan footage of a specific song from a concert, you're searching YouTube for 20 minutes hoping someone uploaded it. On your platform, you tap "Bodak Yellow" on the setlist and get 8 clips of that exact song from that exact show, from different angles, in one place.

The tagging can happen two ways:

- **Manual by uploader** — "what song is playing in this clip?" with setlist pre-populated as options (low friction because you've pulled the setlist from setlist.fm already)
- **AI-assisted** — audio fingerprinting or a simple AI classifier can identify the song from the video's audio automatically, pre-tagging uploads without requiring any effort from the uploader

That second option is the moat builder. If you auto-tag song clips, your platform becomes dramatically more useful than anything else out there, and you've done it without burdening

uploaders at all.

The "vibe meter"

This might sound gimmicky but it's actually a meaningful piece of information for non-attendees. If someone is deciding whether to watch 2 hours of fan footage from a show, a simple energy/crowd/staging score tells them what they're getting into. It can be computed from:

- Like velocity on uploads (high energy = clips getting liked fast)
- Video audio levels (crowd noise)
- Caption sentiment from uploaders
- Upload density per minute of show (the best moments get captured most)

Over time this also becomes a way to compare shows — "which Cardi B show had the highest energy crowd this tour?" That's the kind of question a superfan is genuinely curious about.

What this framing does to your competitive position

The YouTube comparison you made is apt, but there's a sharper way to say it: **YouTube has this content, but it's completely unsorted.** Search "Cardi B Chase Center 2026" on YouTube and you get a random handful of clips uploaded by different people, in no particular order, with no structure. You get five uploads of the same song and nothing of three others.

Your platform takes that same raw material — fan-shot phone footage — and turns it into a structured, navigable show experience. That's the value proposition. You're not competing with YouTube on video hosting. You're competing with the chaos of YouTube by being the organized, event-specific destination.

This also clarifies the brand direction significantly. You're not "fan photo archive." You're something more like **"the show, from everyone who was there"** — a tagline that speaks directly to the person who couldn't afford to go, couldn't get tickets, lives in another country, or is a casual fan who's curious what the show was like without committing to a full concert film.

How this changes monetization too

The virtual attendance angle opens a door that the archive framing doesn't: **you become relevant to people who have never been to a live event in their life.** A 15-year-old in rural Arkansas who loves Cardi B and will never afford a concert ticket is a real user. A dad in another country whose kid is obsessed with a band he's never heard of is a real user. These people have no ability to upload anything — they're pure consumers. And pure consumers are the majority audience on any YouTube-style platform.

That has implications for where the real revenue ceiling is. An archive serves a niche. A virtual attendance platform serves everyone who's ever been priced out of an experience — which is

most of the world for most major events.

Want to dig into how the Watch mode experience would actually work — the sequencing logic, the AI tagging, or the venue map feature?

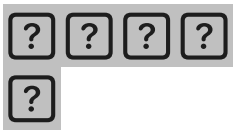

You said: How does setlist-to-video linking work?

How does setlist-to-video linking work?

May 24

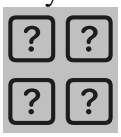

Claude responded: Great question — this is one of the most technically interesting and highest-value features on the platform.

Great question — this is one of the most technically interesting and highest-value features on the platform. Let me break down exactly how it works end to end.

There are two paths to get a video linked to a song: **uploader tags it manually**, or **the platform tags it automatically**. In practice you want both running in parallel, with auto-tagging filling in the gaps that manual tagging misses.

V

visualize

V

visualize show_widget

Path A — Manual tagging

This is the simpler path and it should be the fallback experience for every upload. The key

insight is that you make it **almost zero effort** by doing all the setup work ahead of time.

When a fan uploads a video and attaches it to a Cardi B — Chase Center — March 22 event, the upload UI shows a single optional question: "*What song is playing in this clip?*" Below it is a pre-populated list pulled from setlist.fm for that exact show — not a blank text field, not a search bar, just a scrollable list of the 18 songs she played that night in order. The fan taps the right one and they're done.

A few things make this work well in practice:

Order matters. Show the setlist in performance order, not alphabetical. If someone is uploading a clip from the end of the show, they're scrolling to the bottom — that's correct. Don't make them hunt.

"I'm not sure" is a valid answer. Don't force a tag. An untagged video is better than a wrongly tagged one. The auto-tagging pipeline can catch it later.

Multi-song clips. Sometimes a 3-minute video spans two songs. Let uploaders tag a primary song and optionally a secondary. Or just tag the song that starts the clip.

Post-upload editing. Fans can go back and add or correct a tag after uploading. This is important for the first few hours after a show when the setlist might not even be in setlist.fm yet — fans upload, then tag once the setlist is confirmed.

Path B — Automated audio fingerprinting

This is where it gets technically interesting. Audio fingerprinting is a mature technology — it's how Shazam works, how SoundHound works, how every "what song is this?" feature works. For your use case, the two best API options are:

ACRCloud — the industry standard for this. Used by Spotify, TikTok, and hundreds of other platforms. Recognizes songs from partial audio clips (as short as 10 seconds), works even with crowd noise layered on top, returns song title, artist, and timestamp offset within the song. Pricing is usage-based and very affordable at startup scale — roughly $0.001 per recognition request.

AudD — similar capability, slightly simpler API, good documentation, also handles humming/singing recognition which you don't need but it shows the robustness of the engine.

Here's exactly what happens when a video is uploaded:

1. **Audio strip** — on ingest, your processing pipeline extracts the audio track from the video file. This is a trivial ffmpeg operation, takes milliseconds.
2. **Fingerprint sample** — you don't need to send the entire audio file. Send 20-30 seconds from the middle of the clip (skipping the first few seconds which are often crowd noise or the fan moving their phone). The API returns a match with a confidence score.
3. **Confidence threshold** — if the match comes back at 85%+ confidence, you auto-tag the

video to that song. If it's below that threshold, you flag it for the community review queue. If it returns no match at all (happens when the fan is filming during a quiet moment, a speech, a gap between songs), you leave it untagged and serve it in the "untagged clips" bucket.

4. **Cross-check against setlist** — here's the important refinement. You're not just trusting the fingerprint match blindly. You cross-reference the returned song against the known setlist for that event. If the fingerprint says "Bodak Yellow" and Bodak Yellow is song #3 on the setlist for this show, you tag it with high confidence. If the fingerprint returns a song that wasn't on the setlist, you flag it for review — either the fingerprint is wrong, or the artist played an unplanned song that setlist.fm doesn't have yet.
5. **Timestamp offset bonus** — ACRCloud also returns how far into the song the clip starts. If Bodak Yellow is 3 minutes 42 seconds long and your clip matches at the 2:10 mark, you know the fan captured the last 90 seconds of that song. This lets you eventually sort clips by "which part of the song" — opening bars vs. climax vs. ending — which is genuinely useful for virtual attendees trying to see the peak moment of each song.

The community correction layer

Auto-tagging will be wrong sometimes. Crowd noise, bad audio quality, a cover song, an unusual remix — these can all confuse the fingerprinter. So you need a lightweight community correction mechanism:

Dispute button on each tagged video — "this tag looks wrong." Three disputes from different users triggers a review flag. The correction is surfaced to the uploader ("did you film this during WAP or Bodak Yellow?") and to any contributor with enough trust history to vote.

Confidence display — optionally show users the confidence score on auto-tagged videos, like a subtle indicator. High-confidence tags need no action. Low-confidence tags invite correction.

Uploader always wins — if the person who filmed the video manually tags it, that overrides any auto-tag. They were there. They know.

What this unlocks for the viewer

Once you have setlist-linked videos at scale, a few genuinely powerful experiences become possible:

V

visualize

V

visualize show_widget

This is what the viewer sees when they tap a song on the setlist. A few things worth noting in the design:

The show timeline bar at the top shows where in the overall show this song fell. For a virtual attendee, this is orienting — it tells them they're 18% into the show, not just floating in abstract clip space.

Sort modes are where the real UX power is. "Best angle" is the default algorithmic pick. "Floor / pit" and "Upper sections" let you choose your vantage point preference — just like choosing your ticket. "Full song only" filters to clips where the uploader captured the entire song, not just a fragment. These are the kind of controls that make someone feel like they're genuinely navigating the show.

Confidence scores are visible but small — they're for trust, not for feature. A 97% confidence auto-tag next to a manually tagged clip tells you both can be trusted. The 71% confidence clip is visually dimmed and flagged — the viewer knows to take it with a grain of salt.

Song navigation at the bottom is essential. A virtual attendee is going to step through the whole setlist song by song, like chapters in a movie. Prev/next makes that feel intentional, not accidental.

The edge cases worth solving early

Songs not on setlist.fm yet. In the hours right after a show, the setlist might not be entered yet — fans do this manually on setlist.fm, usually within a day. In the meantime, clips get tagged as "unconfirmed song" and are surfaced in a holding bucket. Once the setlist is confirmed and imported, all those clips get matched retroactively.

Surprise songs / covers. An artist plays a cover or a song not in their normal catalog. The fingerprint will still identify it (ACRCloud knows virtually every commercially released song), it just won't match anything on the imported setlist. These get flagged as "unlisted song" — which is actually a feature, not a bug. Superfans love discovering that an artist played a rare cover, and your platform surfaces it automatically.

Instrumental sections, crowd chants, speeches. These don't fingerprint to a song. They go into an "untagged moments" bucket that can be browsed separately. Actually useful — sometimes the between-song banter is more interesting than the songs.

Live versions that sound different. Some artists remix songs heavily live — the live version might not fingerprint well against the studio recording. This is the hardest case. Mitigation: use timestamp cross-referencing. If 8 clips from this show all fingerprinted as "no match" but were uploaded within a 5-minute window, and the setlist shows a song at that time, you can infer they're probably from that song and offer that as a suggested tag to uploaders.

Why this is a real technical moat

Here's the thing — the fingerprinting API itself is a commodity. Anyone can call ACRCloud. The moat isn't the technology, it's the **combination of setlist data + fingerprinting + a crowd of uploaders tagging things + a correction layer + time.** After two years of shows, you'll have a dataset of confirmed song-to-clip links that no one else has. That dataset trains your own confidence models, reduces your API costs as your local cache fills up, and becomes genuinely defensible.

It also compounds in an unexpected way: once you have a critical mass of tagged clips per song, you can start answering questions no one else can. Which song in Taylor Swift's setlist generates the most fan video? Which moment of a show gets filmed the most? Which seat sections capture the best angles for specific songs? That's data that artists, tour managers, and venue operators would genuinely pay for — and it falls out naturally from the tagging system you've already built.

Want to go into the venue map feature, the Watch mode sequencing logic, or something else?

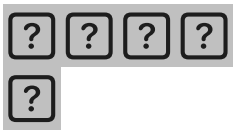

You said: love it.

love it. how should i approach building the MVP?

May 24

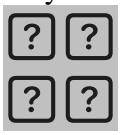

Claude responded: Great question — and the answer is probably more constrained than you might expect.

Great question — and the answer is probably more constrained than you might expect. The goal of an MVP isn't to build a simplified version of everything. It's to find the **one core experience** that proves the concept with real users, and build only what's needed to deliver that.

Let me lay out exactly how I'd approach it.

The core question first: what are you actually validating?

There are two different bets embedded in this product and they require different MVPs:

Bet 1 — Supply: Will fans actually upload their media to an event page? (Upload behavior)

Bet 2 — Demand: Will non-attendees actually spend time watching fan footage to "virtually attend" an event? (Consumption behavior)

Most products fail because they try to validate both simultaneously and end up with a half-built version of each. The smarter move is to sequence them — validate supply first, because without uploads there's nothing to consume. But you want to do it in a way that also gives you early signal on demand.

What the MVP actually is

V

visualize

V

visualize show_widget

The MVP is deliberately brutal in its scope. Here's the reasoning behind every cut:

Why no app yet? A web upload experience on mobile is good enough to validate upload behavior. React Native adds 2-3 months to your build time and you don't need it to answer the core question. Build the web experience mobile-first (responsive, fast, camera roll accessible via browser file picker) and ship.

Why no audio fingerprinting yet? Manual song tagging by uploaders is sufficient to prove that fans will tag their clips. You don't need auto-tagging to validate the concept — you need it to scale. Adding ACRCloud before you have upload volume is over-engineering.

Why no user accounts? Accounts are friction. Every required account creation step costs you roughly 60-70% of potential uploaders who would have otherwise uploaded. Anonymous uploads with an optional "add your email to track your uploads" is enough for the MVP. Accounts come when you need to build uploader identity and community features.

Why manual setlist entry? The Ticketmaster and setlist.fm API integrations are real engineering work. For your first 10-20 events, you create the event pages and type in the setlist yourself. This takes 5 minutes per show. It also keeps you in the loop on what's working before you automate it.

The build plan: 8 weeks

V

visualize

V

visualize show_widget

The stack decision

You have two realistic paths depending on your technical background and budget:

If you're building yourself or with a small team:

- **Framework:** Next.js 14 (App Router) — SSR is non-negotiable for SEO, and Next.js is the current standard
- **Database:** Supabase — gives you Postgres, auth (for when you need it), storage, and a decent admin UI all in one. Free tier is generous.
- **Media storage:** Cloudflare R2 for photos (free egress is the killer feature), Mux for video
- **Deployment:** Vercel for the web app, dead simple Next.js deploys
- **Search:** Typesense Cloud when you need it (cheaper than Algolia, self-hostable option)
- **Total cost at MVP stage:** Under $50/month

If you're hiring a dev or agency:

Same stack, but you can move faster by using a service like Uploadcare or Cloudinary for the upload + media pipeline instead of wiring R2 + Mux manually. Costs more per GB at scale but saves 2-3 weeks of infra work.

The launch strategy — don't launch to everyone

The biggest mistake first-time founders make with a platform like this is posting on Product Hunt on day one and then watching the tumbleweed. You need content before you need visitors. Here's the sequence:

Step 1 — Pick one artist, one show, one city (weeks 1-6, while building)

Don't try to cover all entity types in the MVP. Pick music concerts first — they have the highest upload motivation (fans *want* to share their concert experience), the most passionate communities, and the most available setlist data. Pick one artist with a tour coming up in the next 4-6 weeks. Not a stadium act. A touring band or artist with 50k-500k followers, passionate fans, and a tour hitting multiple cities. Think: a popular indie artist, a mid-tier hip-hop act, a rising country artist — someone whose fans are deeply engaged but who doesn't have a legal team watching for IP violations.

Step 2 — Seed the first event before anyone else uploads

Before you tell a single person about the platform, attend or recruit someone to attend the first show and upload 20-30 pieces of real content yourself. A mix of photos and short videos, from different seat locations. This seeding is essential — an empty page gets zero organic uploads. A page with 25 existing uploads gets more uploads, because people can see it's real and active.

Step 3 — Tell the artist's fan community, not the artist

Don't pitch the artist yet. Find their fan subreddit, their fan Discord, their fan Twitter/X accounts. Post something like: *"I built a place to upload and watch fan footage from [Artist] shows — here's the page for [Show] in [City], just happened last week. Who was there?"* You want the community response before you're on the artist's radar.

Step 4 — Measure before you build more

After 2-3 events with real upload behavior, you'll know: how many people upload per show (your supply metric), how many people view without uploading (your demand metric), and which features people are asking for in comments/DMs. That data tells you what to build in Phase 2 — not your assumptions.

The metrics that actually matter at MVP stage

Not DAU. Not revenue. These three:

Uploads per event — your core supply health metric. Even 10-15 real uploads from a single show proves the behavior is real. 50+ from a single show is genuinely exciting. This tells you whether the upload motivation is there.

View-to-upload ratio — for every person who uploads, how many people just watch? A healthy ratio is 10:1 or higher. If you're at 100 uploaders and 2,000 viewers on an event, you're building something. This validates the demand side.

Return visits — does anyone come back to the same event page more than once? This is the strongest signal that the virtual attendance concept is working. Someone who watches clips from a show, goes away, and comes back the next day to watch more is exactly the behavior you're trying to create.

The one thing that will make or break the MVP

Everything above is table stakes. The actual make-or-break factor is this: **the upload experience on a phone, 5 minutes after a show ends, has to be under 60 seconds from opening the site to finishing an upload.**

Think about the context. The show just ended. People are filing out. They're emotional, excited, maybe a little deaf from the speakers. They have 15 minutes before they're in an Uber or on the

subway losing signal. If uploading to your platform takes longer than posting to Instagram Stories, they will post to Instagram Stories instead and you get nothing.

That means:

- The event page URL needs to be reachable from a QR code they scanned on the way in (or a link in bio they tapped before the show)
- No forced account creation before uploading
- The file picker opens immediately — no multi-screen onboarding flow
- Upload starts in the background the moment they confirm — they can close the browser and it still completes
- The only question you ask is the song tag — and it's optional

If you get that right, everything else is secondary. The platform lives or dies on whether fans actually upload, and that decision happens in a 60-second window right after the show.

Want to go deeper on the tech stack, the SEO strategy, or how to approach that first artist community launch?

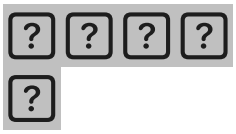

You said: do you think that apps like CrowdAlbum, sureshot, Happaning etc have 'proven' that users are willing to do the upload part?

do you think that apps like CrowdAlbum, sureshot, Happaning etc have 'proven' that users are willing to do the upload part? feels like theres some insights from existing companies/apps we can use to make reasonable assumptions

May 24

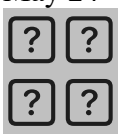

Claude responded: Really sharp instinct — yes, absolutely.

Really sharp instinct — yes, absolutely. You don't need to validate from zero. Let me research what actually happened with these companies and what we can extract from their trajectories.

Searched the web

Searched the web

Good. That's a rich set of data to work with. Here's the honest synthesis.

Yes — and here's exactly what's been proven versus what hasn't

The short answer is: **upload behavior is well-proven. Virtual attendance consumption is not.** That's actually great news for your MVP strategy — it means you're de-risking the harder question, not the easier one.

Let me break it down company by company, then pull the underlying data.

What CrowdAlbum actually proved

CrowdAlbum worked with more than 1,000 artists and venue partners before being acquired — on only $100k raised and a tiny team. That's remarkable distribution for a bootstrapped startup. But here's the crucial detail: CrowdAlbum gathered photos and videos shared on Twitter and Instagram without fans needing to use hashtags, presumably using time and location data from social media posts. TechCrunchFortune

Read that carefully. CrowdAlbum wasn't asking fans to upload *to* their platform. It was **aggregating content fans had already posted elsewhere**. The upload behavior it proved was "fans post concert content to Instagram and Twitter" — which is undeniably true, but it's not the same as "fans will upload directly to a dedicated platform."

That's an important distinction. It means CrowdAlbum sidstepped the cold-start upload problem entirely — and also means it never actually tested whether fans would upload natively. Spotify bought it for the artist relationship tools and the data layer, not for a proven direct-upload product.

What it does prove: Fan-generated concert content exists at massive scale on social platforms. The behavior is real. The content is there. The question is whether fans will redirect that behavior to you.

What SureShot and Happaning prove

These are B2B products — SureShot is a Copenhagen-based startup building an AI-assisted crowd-generated video platform for event organizers. Organizers pay, set up an event, distribute a PIN, and collect fan uploads. They've proven that **when an organizer actively promotes a collection campaign at their event — signage, announcements, incentives — fans will upload directly to a platform.** SureShot

That's a stronger proof point than CrowdAlbum. But there's a catch: the organizer is doing real work to drive participation. Signage at the venue, announcements from the stage, sometimes prizes for best clip. The upload rate when an organizer is actively pushing it vs. organic

discovery is likely 5-10x different.

Your platform is more organic — you're not relying on an organizer to push it. That's harder, but it's also how you scale without a B2B salesforce.

The underlying behavior data is actually very strong

This is where it gets genuinely encouraging. The data on fan filming and sharing behavior is extensive:

56% of concert attendees engage on social media at a show — including 74% of millennials — with 33% sharing photos and 25% posting videos. That's Nielsen data from a large sample. At a 20,000-seat arena, that's roughly 5,000 people posting videos and 6,600 sharing photos during the show itself. Billboard

Research on fan filming behavior found that fans selectively upload clips with the highest excitement — songs with lots of screaming, group behaviors like chanting and cheering in unison. The clips often included songs with special meaning to the fandom. This is crucial — fans aren't filming randomly. They're filming the moments they want to relive and share. That's exactly the content you want on your platform. Sage Journals

On the demand side — the people who couldn't attend: a Live Nation global study of 40,000 people found fans traveled 40 billion miles for concerts in 2024, with 93% saying they attend shows because they crave real experiences over virtual ones. That last number sounds like bad news for virtual attendance, but read it again — 93% of people *who attended* prefer real. The question is about the much larger group who never got there. marketscreener

And that group is growing. The average ticket price for a concert in one of the top 100 global tours was $136 in 2024, up 50% from an average of $91 in 2019. One in five Gen Z concertgoers reported spending beyond their means to attend concerts in the past two years. The pricing pressure is real and structural — it's not a temporary blip. More people are being priced out every year, and they're not losing interest in the music. CNBCBillboard

The one genuinely unproven thing

Here's the honest gap in the validation stack:

Nobody has proven that a non-attendee will spend 30+ minutes watching fan-shot concert footage on a dedicated platform.

YouTube proves people watch concert footage — but usually officially released content, or occasional viral clips. The closest analog is the "fancam" culture on Twitter/X and TikTok, where individual clips from concerts rack up millions of views. That proves demand for clips. But a sustained "virtual attendance" session — working through a setlist, watching multiple angles, spending real time with an event — that's an unproven behavior at platform scale.

This is actually fine. It's a reasonable inference from strong adjacent signals, and you don't need to prove it before building the MVP. But it's the question your MVP should be designed to answer, not assumed.

So what does this mean for your assumptions and your build?

V

visualize

V

visualize show_widget

The practical takeaway for your build

This validation map tells you something important about sequencing:

You don't need to prove upload behavior from scratch — the adjacent evidence is strong enough to proceed with reasonable confidence. Fans film everything. They share everything. They contribute to archives they care about. The specific question of "will they upload directly to your platform without organizer prompting" is the one gap, and the answer is almost certainly yes *if* the friction is low enough and the page they're uploading to has existing content. Which is why the seeding strategy from the MVP plan matters so much.

The virtual attendance session length is the only thing you genuinely don't know. And here's the good news: you don't need to solve it in the MVP. You just need to measure it. Put Plausible or Fathom analytics on the event pages from day one, track average session duration and page depth, and within 3-4 events you'll have real data. If non-attendees are spending 8 minutes on an event page, you have a signal. If they're spending 45 minutes, you have a product.

The pricing-out trend is your structural tailwind. Concert ticket prices are up 50% since 2019 and every industry analyst expects them to keep rising. One in five Gen Z fans already spent beyond their means to attend shows. The Live Nation CEO has literally said tickets are still underpriced. This isn't a temporary condition — it's the permanent reality of the live events industry. Every year that passes, your potential audience of priced-out fans gets larger. That's a tailwind most startups would pay for. CNBCBillboard

The companies before you proved the content exists and the desire exists. You're the first to try to organize it into a real viewing experience for the people who weren't there. That's the unproven piece — and it's also the most valuable one if it works.

You said: can i seed some of the entity pages by scraping existing social content to fill the page with, perhaps?

can i seed some of the entity pages by scraping existing social content to fill the page with, perhaps?

May 24

Claude responded: Great instinct — and this is exactly what CrowdAlbum did to bootstrap.

Great instinct — and this is exactly what CrowdAlbum did to bootstrap. But it's a legal and technical minefield you need to navigate carefully. Let me research the current state of scraping law before giving you a real answer.

Searched the web

Searched the web

Good — now I have a clear picture. Here's the honest, nuanced answer.

The short answer

You can seed pages with social content, but **how** you do it matters enormously. There are three distinct approaches with very different legal profiles, and you should only pursue two of them.

The three approaches, ranked by risk

1. Embed, don't host — low risk ✓

This is the approach that's cleanest legally and is also how most of the web works. Instead of downloading content and storing it on your servers, you embed it directly from the source platform using their official embed API.

Instagram, TikTok, YouTube, and Twitter/X all provide official oEmbed endpoints. When you embed a post this way:

- The media file never touches your servers
- The original creator's username and link are displayed automatically
- The platform's own infrastructure serves the content
- Instagram's own API documentation explicitly describes the oEmbed endpoint and states "you can get a public Instagram photo or video post's embed HTML code and use it to embed the post in other websites — each embedded post gives proper attribution by displaying the post owner's username and linking to the post owner's profile." NatLawReview

The legal protection here comes from what courts call the "server test" — infringement only occurs if a copy of the image is stored on the accused infringer's computer server or other storage device. Embedding via official APIs passes this test because you're not storing anything. Miller Nash LLP

There is some ongoing legal murkiness — a photographer petitioned the Supreme Court in 2025 to reconsider the server test — and there's a circuit split, with the Second Circuit saying embedding can infringe and the Ninth saying it doesn't. But for a startup seeding content in the short term, embedding via official APIs is the pragmatic, defensible choice. Just be ready to pull embeds if a platform changes its policy or a creator requests removal. Digital Camera WorldPassle

What this looks like in practice: You search Instagram for `#beyonceconcert` `#renaissancetour` and find 50 good fan posts. You pull the oEmbed code for each, attach them to the correct event page, and they display with full attribution. The fan sees their post featured on your platform, often likes it, and sometimes becomes an uploader themselves. That last part is actually a flywheel.

2. Aggregate public metadata + link out — very low risk ✓

Even safer than embedding: don't display the content at all on your page. Instead, surface a feed of public posts with thumbnails and captions that link out to the original. Think of it as a curated directory, not a media host.

"12 fans posted about this show on Instagram — see them here →"

This is essentially what Google Image Search does. The hiQ v. LinkedIn ruling established that scraping publicly available data on the web doesn't violate the Computer Fraud and Abuse Act, and the 2024 Meta v. Bright Data ruling reaffirmed that scraping publicly available data remains legally defensible — Meta lost because they couldn't prove Bright Data was scraping data behind login walls. ScrapflyScrapecreators

Linking out drives traffic to the original creator, which is genuinely good for them, builds goodwill, and keeps you firmly in the clear. The downside is a weaker on-page experience — someone has to click away to see the content.

This is also a good long-term partnership feature. "Your post was featured on [Platform] as a highlight from the Beyoncé show — come upload your originals directly." That notification, sent

via Instagram DM or tagged in a comment, converts social posters into direct uploaders.

3. Scrape and re-host — high risk ✗

Downloading images and videos from social platforms and storing them on your own servers is where you get into real trouble, for two separate reasons:

Copyright: The fan owns their photo or video. You don't have a license to copy and store it. The server test no longer protects you once the file is on your server. Even if the content is publicly visible, copying it is a different legal act than linking to it.

Platform ToS: Every major platform — Instagram, TikTok, Twitter/X — explicitly prohibits scraping content for re-hosting in their terms of service. LinkedIn's court victory over hiQ ultimately came down to contract-based ToS violations, not just the CFAA. Platforms have gotten more aggressive enforcing this, not less. Social Media Today

The Clearview AI cautionary tale: Clearview AI scraped billions of images from social media and in 2025 settled a class action lawsuit agreeing to pay approximately $51 million and award plaintiffs 23% of company equity. Obviously facial recognition is a different context, but it illustrates what scrape-and-store at scale looks like from a liability perspective. Lection

Don't do this. The risk-to-reward ratio is terrible for a startup that's trying to build partnerships with artists and venues.

The smarter seeding strategy: a combination approach

Here's how to think about it as a sequenced playbook:

V

visualize

V

visualize show_widget

The flywheel this creates

The notification play in Phase 1 is worth emphasizing because it's not just a legal workaround — it's actually a better growth strategy than silent scraping. Here's why:

When you comment "we featured your clip on [Platform]," several things happen. The creator sees they have an audience beyond their own followers. Many will be genuinely flattered. Some

will upload their full-resolution original directly to your platform, which is far better content than a compressed social media re-post. And a percentage will mention the platform to their followers — "someone built this cool archive of the show, my video is on it."

Silent scraping gets you content. Notifying creators gets you content *and* word-of-mouth *and* a conversion path to native uploaders. It's strictly better.

The one practical wrinkle: API access

The embed approach depends on platform APIs not breaking. A few things to know:

Instagram/Meta is the trickiest. Meta has been restricting API access progressively since Cambridge Analytica. The oEmbed endpoint still works for public posts, but the ability to search for posts programmatically is limited without app review. For Phase 1 manual seeding, this doesn't matter — you're finding posts manually and embedding them. For Phase 2 automation, you'd need to apply for Meta's developer program.

TikTok Research API requires an application and approval, but it exists and is legitimate. Worth applying early.

Twitter/X has a free Basic API tier that gives you limited search access. Sufficient for discovery at MVP scale.

YouTube Data API v3 is generous on the free tier — 10,000 units per day, which is plenty for searching concert footage.

Bottom line

Yes, you can and should seed pages with social content — but do it as **embeds with attribution, not downloads with re-hosting**. The legal risk of embedding via official APIs is low and manageable. The legal risk of scraping and storing is real and not worth it for a startup that's trying to build industry relationships.

The seeding strategy also doubles as your first growth loop: find the fans who are already posting, feature their content, tell them about it, convert them to native uploaders. That's a better outcome than any amount of silent scraping.

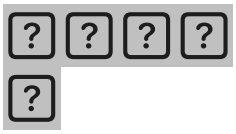

You said: write a two-pager overview pdf of this idea.

write a two-pager overview pdf of this idea. page 1 is the overview, page 2 is the visual mockup

you already made for the cardi b page (change it to Morgan Wallen)

May 24

Claude responded: Let me read the PDF skill first, then build this out.

Let me read the PDF skill first, then build this out.

Reading PDF skill

Reading PDF skill

Good. Using reportlab. Let me build this now.

Ran 2 commands, created a file, read a file

Ran 2 commands, created a file, read a file

There it is — two pages, ready to share. Here's what's in it:

Page 1 covers the full concept: the problem (ticket prices, priced-out audience), the solution (Showside as virtual attendance), how it works in three columns, the competitive gap showing where CrowdAlbum/SureShot fell short, and four market signal bullets with the key stats we surfaced in the research.

Page 2 is the Morgan Wallen entity page mockup — browser chrome, nav bar, hero with stats, the tabbed event list with upcoming and past shows, and the "most clipped songs" strip at the bottom showing the setlist-linking concept.

Let me know if you want to adjust anything — name, colors, copy, layout, or add a third page.

Showside overview
Document · PDF

Open in Google Chrome

You said: sorta like the 'shorts' section on youtube but way better.

sorta like the 'shorts' section on youtube but way better...

May 24

Claude responded: Yes — exactly.

Yes — exactly. And this screenshot actually makes the point better than anything we've said so far. Look at what YouTube is showing you for "taylor swift live":

- Cruel Summer LIVE at the Eras Tour 5-27... **11M views**
- Taylor's disappointed face when they started to sing early... **16M views**
- When one surprise guest brings out the ULTIMATE... **13M views**

These are fan-shot phone videos. Millions of people are watching them. The demand is **already proven and massive** — it's just completely unorganized on YouTube. No one knows which show that Cruel Summer clip is from. You can't find all the clips from that specific night. You can't step through the setlist. You can't see what the show looked like from section 200. You just get an algorithm-sorted pile of whatever YouTube decides to surface.

That's the gap in one screenshot.

What YouTube gets right (and you inherit)

- Fans already upload concert footage compulsively — the content supply is real
- Non-attendees already watch it in huge numbers — the demand is real

- Short vertical clips (Shorts format) perform especially well — your mobile-first upload experience should lean into this
- People discover content through search, not just feeds — SEO is a real acquisition channel

What YouTube gets wrong (and you fix)

YouTube	Showside
Search returns a random mix across all dates and venues	Every clip attached to a specific show
No way to find "all clips from Eras Tour Philadelphia night 2"	Exact event page with every upload from that night
Algorithm buries older clips	Permanent archive — 2019 show is as findable as last week
No setlist context — you don't know what song is playing	Clips tagged to songs, browsable by setlist position
Single creator's channel owns the content	Community archive — 200 uploaders contributing to one page
No seat-section filtering	Browse by floor, section 100s, upper deck
Shorts and long videos mixed with no event structure	Structured by event → song → angle

The one thing to add to your pitch

This screenshot is your slide 2 in any investor or partner conversation. Slide 1: "Ticket prices are up 50% and millions of fans are priced out." Slide 2: this screenshot with the caption "**The demand is already there. Nobody has organized it.**"

YouTube is accidentally proving your thesis every day. Millions of views on fan-shot concert footage, zero structure around it. You're not building something people need to be convinced to want — you're building the organized version of something they're already consuming.

Want to update the PDF to add this angle, or keep going on the build plan?

?????

You said: update the pdf, add more details about the idea and how it will work, you can make it 3 pages long, and also make sure the design is cleaned up.

??

update the pdf, add more details about the idea and how it will work, you can make it 3 pages long, and also make sure the design is cleaned up. use these image of wallen to bring the pdf to life more

May 24

Claude responded: Let me look at the current PDF output and plan the upgrade properly.

Let me look at the current PDF output and plan the upgrade properly.

Ran 3 commands

Ran 3 commands

All images converted. Now building the full 3-page PDF.

Ran a command, read a file

Ran a command, read a file

Here's the updated 3-pager. Here's what changed across each page:

Page 1 — The Opportunity. Full-bleed hero image (the B&W crowd/stage shot) with a dark gradient overlay and the headline "*The show, from everyone who was there*." The problem, insight, and solution sections are fully written out, with the three-column how-it-works breakdown at the bottom.

Page 2 — Product & Market. Opens with a two-image strip (stadium shot + bokeh lights) to anchor the visual feel. Left/right column breakdown of the upload experience vs. the watch experience, the full competitive table, and the four market signal bullets with all the stats we pulled from research.

Page 3 — The Morgan Wallen mockup. Cleaned-up browser chrome, real photo integrated into the hero band (guitar photo fading into the dark background), stats strip, tabbed event list with upcoming and past shows, most-clipped songs, and a four-photo strip at the bottom using your uploaded images as a closing visual.

Let me know if you want any copy tweaks, layout adjustments, or a different photo placement.

Showside overview v2
Document · PDF

Open in Google Chrome

May 24

Claude responded: Good research instinct — you went and checked yourself.

Good research instinct — you went and checked yourself. Let me look at these properly and then research each one.

Searched the web

Searched the web

Good research. This is really useful — and genuinely more competitive than what we found in our initial research. Here's the honest, full picture.

What these screenshots are showing you

You Googled your own concept and Google's AI surfaced a more complete competitive landscape than we found initially. That's healthy — better to know now. Let me go through each one specifically.

Concert Archives — the closest existing thing

Concert Archives is a social network and diary for music fans to remember and share concerts they've attended. Under each specific show date and venue, fans can upload photos and videos, view media uploaded by other attendees, and check the setlist. It's been running since 2013, has been featured in The New York Times, Newsweek, The Guardian, Reuters, BuzzFeed, and Pitchfork. SoftonicGoogle Play

This is genuinely the closest thing to what you're building. It has the entity → event hierarchy, it has photo/video uploads, it has setlist integration. You need to take it seriously.

But here's what the reviews tell you: Users consistently say "great idea, needs a LOT of work" — the app overheats phones, archives don't load, there are duplicate shows, bugs crash sessions after hours of data entry. It's a passion project that never got proper product investment. Their Facebook page has 1,560 likes. After 12 years. That's not a thriving platform — that's a concept that validated demand but never had the execution to capture it. JustUseAppFacebook

The critical gap: Concert Archives is built for attendees logging their own history. It's a personal diary with sharing features. It is NOT built around the virtual attendance experience — the non-attendee consuming the full show. That reframe is still unoccupied.

FanVids — very recent, very relevant

This is the most important new competitor. FanVids launched in March 2026 — a free app that lets artists collect fan videos at live shows, share them with other fans, and collect fan data. Growing since its April 2024 launch, FanVids has worked with hundreds of artists including Kelsea Ballerini and Billy Idol. Monthly users increased tenfold in just three months. CelebrityAccessLimitless-magazine

This is real traction and recent. You need to look at this seriously.

But the model is still B2B-artist-first: FanVids is built for artists, labels, and managers to collect fan UGC and first-party data for their CRM, ticket sales, and merch. The pitch is: "fans upload content directly to you, and you capture their contact info in the process." The fan experience is secondary — it's a marketing data pipeline that also happens to show fans each other's videos. FanvidsRostr

The artist owns the page. The artist controls it. A fan browsing FanVids is browsing the *artist's* marketing tool. That's fundamentally different from a community-owned fan archive where no single party controls the content.

Aftershow.io — same B2B lane

Aftershow.io lets artists collect unlimited videos and photos from fans with smart upload

management, branded URLs, separate landing pages for every concert, geographic visualization of upload activity, and dynamic QR codes for printing on merch and posters. aftershow

Again — artist-controlled, B2B model, fan data capture is the core value prop. Same lane as FanVids and SureShot.

The honest updated competitive picture

Here's what this research actually tells you — and it's more nuanced than "the gap is wide open":

The **B2B/artist-controlled** lane is getting crowded fast. FanVids, Aftershow, SureShot, Happaning — all launching or growing in 2024-2026, all targeting artists and managers as the customer, all using fan uploads as a data capture mechanism for the music industry.

The **consumer/fan-first** lane — where the fan is the primary user, the archive belongs to no single party, and non-attendees are the biggest audience — is still genuinely empty. Concert Archives tried and never had the product quality or the virtual attendance framing to capture it.

That distinction is your entire strategic position. Here's how to say it sharply:

FanVids and Aftershow answer the question: "How do artists collect and own fan content?"

Showside answers the question: "How do fans who weren't there experience the show?"

Those are different products serving different primary users. One is a marketing tool. The other is a media experience. The fact that the B2B lane is busy actually validates that upload behavior is proven at scale — those platforms are getting fans to upload. You just want those same uploads to live in a community archive instead of an artist's CRM.

What you should do right now

Two things:

1. Spend an hour on concertarchives.org — actually use it, upload something, browse a show. Map exactly where it breaks down. That's your product benchmark. Your MVP needs to be noticeably better on execution even if the concept overlaps.

2. Watch FanVids closely — they're 2 years old, growing fast, and artist-first. The risk is that a well-funded artist management company uses FanVids as infrastructure and accidentally builds the consumer experience as a byproduct. That's the scenario where your moat erodes. Getting to consumer scale before that happens matters.

The gap is real but it's narrower than it was 6 months ago. That's not a reason to not build — it's a reason to build faster and be cleaner about the positioning.

You said: conceptually, concertarchives.

conceptually, concertarchives.org is similar - but it looks very dated and doesn't look like users are engaging much

May 24

Claude responded: Yes — and these screenshots tell you everything you need to know.

Yes — and these screenshots tell you everything you need to know. Look at what you're actually seeing:

The Roxy Theatre — 7,475 concerts in the database. The latest video is from **December 2024** and **January 2015**. That's not a typo. The two videos on the page are 1+ year apart. The photo grid for Rockwood Music Hall is almost entirely one person — Pete Blair — uploading 5+ photos from a single 2023 show. 162 attendees listed, almost no media.

This is a database masquerading as a community. The data is there. The engagement isn't.

What Concert Archives actually is vs. what it looks like it should be

It's essentially a **concert attendance ledger** — fans log "I was there" the same way you'd log a book on Goodreads. The media is an afterthought bolted on later. The proof is in the UI: photos and videos are buried under a "Scroll to" nav, the concert list is a plain text table, and the whole thing looks like it was designed in 2013 and never meaningfully updated.

The ads are the tell. Wine ads. Adobe Acrobat ads. YouTube TV ads. These are remnant network ads served to a site that never built a real monetization strategy — which means they never built a real product strategy either.

The five specific things these screenshots show you to do differently

1. Lead with media, not data. Concert Archives leads with a text list of shows. Every page should open with the best video from that venue or artist — autoplay muted, full width. Make someone feel the show in 3 seconds before they read a word.

2. Video is first class, not a tab. On Concert Archives, Videos is a navigation item you scroll to. On Showside, video is the default view. The Watch tab is the hero. Photos are secondary.

3. Upload friction is clearly too high. The sparse media despite 7,475+ concerts at just one venue means people aren't uploading. Either the upload flow is too hard, there's no motivation to do it, or both. Your 60-second mobile upload experience is a direct fix to this.

4. The attendee list is a missed opportunity. 162 attendees at Rockwood but almost no media. Those 162 people were there — they have photos on their phones right now. Concert Archives never connects the logged attendance to an upload prompt. You should: "You logged this show — want to add your photos and video?"

5. No non-attendee experience exists. There is nothing on Concert Archives for someone who wasn't there. It's entirely inward-facing — a diary for people who attended. The entire virtual attendance angle, which is your primary audience, is completely absent. A non-attendee landing on that Roxy page sees a text list of band names and two videos from different years. There's nothing to watch, nothing to experience, nothing to come back for.

What this means for your positioning

Concert Archives accidentally answers a question nobody was really asking — "how do I keep a log of concerts I've been to?" It's Letterboxd for concerts, except Letterboxd actually has a thriving community because movies have a broader audience and logging a film takes 10 seconds.

You're answering a completely different and much larger question: **"I couldn't go to that show — what was it like?"**

That's the reframe that makes everything click. Concert Archives is a diary. Showside is a broadcast. One is private memory management. The other is public entertainment. They happen to share some infrastructure — event pages, media uploads — but the product experience, the primary user, and the value proposition are fundamentally different.

The fact that Concert Archives has been around since 2013, got press in the New York Times and Pitchfork, and still looks like this after 12 years is actually the most bullish signal you could find. The concept has been validated enough to survive for over a decade. The execution has never been good enough to make it matter. That's your opening.

You said: i want to be clear - i definitely like the "I couldn't go to that show — what was it like?

i want to be clear - i definitely like the "I couldn't go to that show — what was it like? angle, but i also see value in the 'archive' angle. i just don't want to lose sight of that as we talk through it. i suppose naturally over time it would automatically become an archive of past events. i do worry a little about trying to capture each individual event and all songs from all those events - it feel like overkill. if a kid in london wants to see a taylor swift concert on our site, she doesn't really care if it was from day 2 in new york or from day 1 in atlanta... ya know what i mean?

May 24

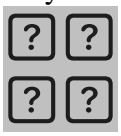

Claude responded: Yes — totally.

Yes — totally. And this is actually a really important product instinct. Let me push on it.

You're identifying a real tension in the architecture

The way we've been talking about it — entity → specific event → setlist → song-by-song clips — is the **archivists's** mental model. Maximum fidelity, maximum organization. Every show catalogued, every song tagged, every date tracked.

But the London kid's mental model is completely different. She wants to answer one question: **"what is a Taylor Swift concert actually like?"** She doesn't care if it was Atlanta night 1 or New York night 2. She wants to feel the energy, see the stage, hear the crowd, watch the moments. She's not researching — she's experiencing.

These are two genuinely different use cases and you're right to hold both.

The way to think about it

Rather than choosing one, think of it as **two views of the same underlying data:**

The experience view — "Taylor Swift live" — surfaces the best content across all shows, all dates, all venues. Crowd energy. Stage production. Iconic moments. Organized by *what it feels like* to be there, not by *which specific night it happened*. This is what the London kid browses. This is the virtual attendance experience. This is also your SEO surface — "Taylor Swift concert

experience" has far more search volume than "Taylor Swift MSG June 14."

The archive view — "Taylor Swift, Madison Square Garden, June 14" — is the specific show page for people who were there and want to find their clips, or for the superfan who knows exactly which night she played a surprise song and wants to find footage of that specific moment. This is Concert Archives territory, but done properly.

The key insight is that **the archive view generates the data, and the experience view surfaces the best of it.** You don't have to choose. The upload always attaches to a specific event — that's how the data gets organized. But the default browsing experience for a casual fan doesn't need to care about that granularity.

What this means for the product

On the entity page — the Taylor Swift page — the default view isn't a chronological list of 152 tour dates. It's a curated **highlights reel**: best videos, most-watched moments, iconic clips, surfaced across all shows. The date and venue are metadata, not the headline. The London kid lands here and immediately sees 20 incredible clips from Taylor Swift concerts. She doesn't need to pick a show first.

The show-specific view still exists — but it's one level deeper. Fans who attended a specific night navigate there to find their show. Superfans who want to compare night 1 vs night 2 go there. The archive is real and valuable. It just isn't the front door.

The setlist-per-song granularity — which we spent a lot of time on — is probably a Phase 2 or power-user feature, not the core. For the London kid, "here are 40 clips of Cruel Summer from various shows" is more useful than "here are 6 clips of Cruel Summer from Atlanta night 1." The song-level tagging is still worth doing because it enables filtering, but the default shouldn't require that level of specificity to have a great experience.

The analogy that makes this click

Think about how Netflix presents a TV show. The show page has a highlights trailer that gives you the feel of the whole series. You don't have to pick Season 2 Episode 7 to decide if you want to watch it. But Season 2 Episode 7 exists and is exactly where you go once you're invested.

Your entity page is the trailer. The specific show pages are the episodes. Most visitors never need to go deeper than the trailer. The people who do are your most engaged users — and they'll love you for having the granularity.

The archive angle doesn't disappear. It becomes the foundation that makes the experience view possible. Every specific upload feeds the machine. The casual viewer just never has to see the machine.

You said: i agree.

i agree. redo the mockup of the entity page with this in mind

May 24

Claude responded: Let me pull up the frontend design skill first.

Let me pull up the frontend design skill first.

Reading frontend design skill

Reading frontend design skill

Good. Now building this with the new mental model — experience-first, archive-second.

V

visualize

V

visualize show_widget

The hierarchy here reflects the new mental model exactly:

What changed from the old version:

The old entity page led with a list of tour dates — essentially a calendar. The London kid had to pick a specific show before seeing anything interesting. That's the archive-first mental model.

This version leads with the feeling. The hero has five actual video clips playing in the background — you see the show before you read a word. Below that, fan highlights are front and center, filterable by mood (crowd energy, acoustic moments, stage moments) rather than by date. You never have to care which show something is from to enjoy it.

The "most filmed songs" section is a useful middle ground — it surfaces the song-level

experience without forcing anyone into a specific setlist. A London kid can tap "7 Summers" and get 241 clips of that song across every show, from every stadium, in one place.

The archive is still there — it's the last element, tucked away as a purposeful "go deeper" door. "142 shows, 2018 to present — browse every show, every setlist, every upload by date and venue." It's real, it's accessible, it just isn't the front door. The people who want that level of detail will find it immediately. The London kid never needs to see it.

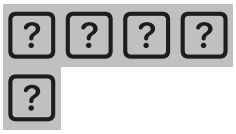

You said: write a complete product scope i can feed into claude code to build me a v1 of the site

write a complete product scope i can feed into claude code to build me a v1 of the site

May 24

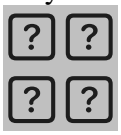

Claude responded: This will be comprehensive.

This will be comprehensive. Let me think through every layer carefully before writing it.

Ran 2 commands, viewed a file, created a file

Ran 2 commands, viewed a file, created a file

691 lines, ~3,800 words. Here's what's in it and why each section is written the way it is:

What you are building — One paragraph that Claude Code reads before anything else. Establishes the mental model so every ambiguous decision gets made in the right direction.

Tech stack — Explicit and non-negotiable. Claude Code will happily invent its own stack if you don't pin it. Supabase + R2 + Mux + Next.js + Vercel is the right combination for this at MVP scale.

Database schema — Full SQL, not a description. Claude Code should be able to run this directly. The schema encodes all the product decisions: anonymous sessions, status flags for moderation, song/section tags, Mux fields separate from R2 fields.

URL structure — Pinned early because every other page spec references it. Stable slugs are called out explicitly because Claude Code will sometimes generate routes with IDs instead.

Page specs — Each page has design intent, not just a feature list. The entity page spec literally says "experience-first, archive-second" and describes the collapsed archive section. This is where your product thinking from this conversation gets encoded into instructions.

Upload flow — The most detailed spec in the document because it's the most important flow. The 60-second rule is stated as a hard constraint, not a preference. Anonymous-first is stated three times in different ways because it's easy to accidentally break.

API routes — Full request/response specs so Claude Code doesn't invent its own. The Mux webhook route is particularly important — without it videos never become viewable.

SEO — Treated as non-negotiable infrastructure, not a feature. The page title templates are given explicitly so they're consistent across all pages.

What is NOT in V1 — Probably the most important section. Without this, Claude Code will start building audio fingerprinting and seat maps.

Acceptance criteria — 10 concrete pass/fail tests. Feed these back to Claude Code at the end to verify everything works before you consider it done.

Showside v1 scope
Document · MD

Open in TextEdit

You said: what model should i use with claude code to build this

what model should i use with claude code to build this

May 24

Claude responded: Good current data.

Searched the web, viewed a file

Searched the web, viewed a file

Good current data. Here's the honest answer:

Use `opusplan` for this project

The `opusplan` model alias gives you an automated hybrid approach: in plan mode it uses Opus for complex reasoning and architecture decisions, then automatically switches to Sonnet for code generation and implementation. For a project of this scope — a full Next.js app with Supabase, Mux, R2, and a dozen interconnected pages — that's exactly what you want. Claude

The scope document you have is essentially the "plan." Opus will reason about how all the pieces fit together. Sonnet will actually write the code fast. You get the best of both without manually switching.

To use it:

bash
```
claude --model opusplan
```

The practical breakdown

Opus excels at initial builds based on project plans, adding entire new modules, complex system design, and multi-system integration. That describes the first session with your scope document perfectly. Steve Kinney

Sonnet 4.6 handles 90%+ of coding tasks without compromise and is preferred over Opus 4.5 by 59% of developers for day-to-day work. Once the architecture is established and you're iterating on individual components, straight Sonnet is fine and meaningfully cheaper. Claude Fast

Max and Team Premium plans default to Opus 4.6, while Pro and Team Standard default to Sonnet 4.6. Claude Code may automatically fall back to Sonnet if you hit a usage threshold with Opus. So check your plan — if you're on Pro, you may not have Opus access at all, in which case Sonnet 4.6 is your answer and it'll handle this build just fine. DEV Community

Practical session strategy

Session 1 — Architecture and database: Use Opus or opusplan. Feed it the full scope document. Let it set up the Next.js project, Supabase schema, and all the API routes before

touching any UI. This is where the reasoning investment pays off.

Sessions 2–4 — Page by page: Sonnet is fine. The architecture is decided, you're just implementing. Start each session with a `/clear` and tell it which page you're building.

Debugging sessions: Switch back to Opus with `/model opus` when something is genuinely broken and you need it to reason across multiple files to find the issue.

One other thing worth knowing: you can run `/model` at any point during a session to switch instantly — you don't have to restart. So start on Opus for the planning phase of a session and drop to Sonnet once you're in pure implementation mode. Claude

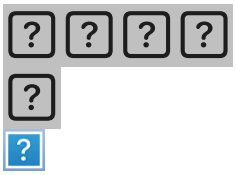

May 24

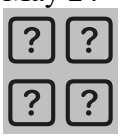

Claude responded: You're already in Claude Code with a Showside project context open — perfect.

You're already in Claude Code with a Showside project context open — perfect. A few things I can see from this screenshot worth noting:

You're on Pro plan — which means Opus 4.7 is available to you (it shows as option 1 in the model picker). That's the newest and most capable model for the architecture/planning work.

You're currently on Sonnet 4.6 · High effort — shown in the bottom right. That's actually a solid default for implementation work.

The Effort setting matters here. For a build this complex, the effort level is arguably as important as the model choice:

- **High** (what you have) — good for most of the build, solid reasoning on each step
- **Max** — use this when you feed it the full scope document for the first time, or when something is genuinely broken across multiple files. It'll think harder and longer before acting.

My recommendation for this specific build:

Start your first session with **Opus 4.7 + Max effort**, paste the entire scope document, and say something like: "*Read this scope document carefully. Before writing any code, tell me your plan for how you'll structure the project, what order you'll build things in, and any questions or*

concerns you have about the spec." Let it reason through the whole thing before touching a file.

Then drop to **Sonnet 4.6 + High** for the actual implementation sessions — that's essentially what you already have set and it'll be faster and cheaper for the mechanical work.

Also noticed you already have a "Showside" project context in the bottom nav — you're set up and ready to go.

You said: claude code results:

```
Plan Project structure showside/ ├── app/ | ├── (marketing)/ # public pages share root layout | | ├── page.tsx # home | | ├── [entity-slug]/page.tsx | | ├── [entity-slug]/[event-slug]/page.tsx | | ├── search/page.tsx | | └── watch/[media-id
```

pasted

claude code results:

May 24

Claude responded: This is a really strong plan — Claude Code has clearly read the spec carefully.

This is a really strong plan — Claude Code has clearly read the spec carefully. Here are my calls on everything it's asking:

Blocking questions

1. Credentials — You'll need to create accounts for all three services before it can test end-to-end. All have free tiers sufficient for development:

- Supabase: free tier at supabase.com
- Mux: free trial at mux.com (no billing needed for low volume dev)
- Cloudflare R2: free tier at cloudflare.com (R2 has no egress fees)

Tell it: "*Create the project assuming I'll supply real credentials. Build with real integrations, not stubs. I'll have all three accounts set up before we start Phase 2.*"

2. Mux webhook in dev — Go with Vercel preview deployment. It's one `vercel deploy` command and you get a real public URL immediately. No ngrok tunnel to manage. Tell it: *"Assume Vercel preview deployment for webhook testing."*

3. "Close this tab" copy — Reword to "leave this tab open." Don't build the server proxy — that's real extra complexity for a minor UX improvement. The honest message is fine.

Design decisions

Returning user sign-in — Option (b), bare `/signin` with magic link not in nav. You'll want this for power uploaders who come back multiple times. Dropping users entirely creates tech debt you'd regret quickly.

Hero clip — Option (a), call it "Featured" and use cumulative view_count. A weekly_views column means a cron job, which is scope creep for V1. Ship it.

Like toggle — Yes, confirm: repeat call = unlike. Standard behavior, correct call.

Stage moments / wide angle — Store aspect ratio from Mux metadata. Mux returns this for free in the asset ready webhook (`aspect_ratio` field). It costs nothing to store and you'll want it later. Don't drop it.

Orphan media rows — Yes, exactly right. `status = 'uploading'` on presign, flip to `active` on complete/webhook. Add a cleanup job note for future but don't build it in V1.

Smaller calls

All fine. Approve everything on that list as stated. The DB triggers for counters is particularly good — that's the right call, not app-level counting.

What to tell Claude Code

Paste this back to it:

Confirmed on all three blocking questions: (1) build with real integrations, I'll supply credentials before Phase 2; (2) assume Vercel preview for Mux webhook; (3) reword to "leave this tab open," no proxy needed. Design calls: (b) magic link signin at /signin not in nav; (a) Featured with cumulative view_count; like toggle = unlike on repeat; store aspect_ratio from Mux metadata. All smaller calls approved as stated. Start Phase 1.

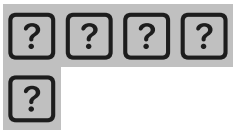

You said: do i need to go create accounts for those platforms now?

do i need to go create accounts for those platforms now?

May 24

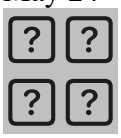

Claude responded: Yes — do it now, before Claude Code hits Phase 2.

Yes — do it now, before Claude Code hits Phase 2. Phase 1 is just scaffolding and database setup, so you have time. Here's exactly what to do for each, in order of how long it takes:

Supabase (5 minutes) Go to supabase.com, create a free account, create a new project. Once it's created, go to Project Settings → API and grab:

- `NEXT_PUBLIC_SUPABASE_URL`
- `NEXT_PUBLIC_SUPABASE_ANON_KEY`
- `SUPABASE_SERVICE_ROLE_KEY`

Cloudflare R2 (10 minutes) Go to cloudflare.com, create a free account. Then in the dashboard go to R2 → Create bucket, name it `showside-media`. Then go to R2 → Manage R2 API Tokens → Create API Token (give it read/write on your bucket). Grab:

- `R2_ACCOUNT_ID` (shown on the R2 overview page)
- `R2_ACCESS_KEY_ID`
- `R2_SECRET_ACCESS_KEY`
- `R2_BUCKET_NAME` = `showside-media`

For `R2_PUBLIC_URL` — in the bucket settings, enable "Public access" and copy the public URL it gives you. You can also set a custom domain later but the default one works fine for V1.

Mux (10 minutes) Go to mux.com, create a free account (no credit card needed for development). Go to Settings → Access Tokens → Generate new token, give it Mux Video read/write access. Grab:

- `MUX_TOKEN_ID`
- `MUX_TOKEN_SECRET`

For `MUX_WEBHOOK_SECRET` — go to Settings → Webhooks → Add webhook. You won't

have a URL yet (that comes when you deploy to Vercel), so enter a placeholder like `https://placeholder.com/api/mux/webhook` for now. It'll generate a signing secret you can copy. You'll update the URL after first Vercel deploy.

Vercel (5 minutes) Go to vercel.com, connect your GitHub account. You don't need to deploy anything yet — just have the account ready. When Claude Code finishes Phase 1 and pushes to GitHub, you'll import the repo into Vercel and add all the env vars in one shot.

Once you have all the credentials, create a `.env.local` file in the project root and paste them all in. Then tell Claude Code the credentials are ready and to proceed to Phase 2.

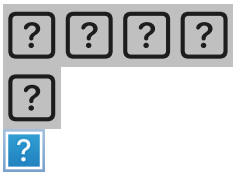

May 24

Claude responded: A few things to fix before you hit "Create new project":

A few things to fix before you hit "Create new project":

1. Project name — change it from "roombox-stack's Project" to `showside`. You don't want this mixed up with your RoomBox stack.

2. "Automatically expose new tables" — uncheck this. Supabase is literally telling you to disable it. Since the spec has explicit RLS policies, you want manual control over what's exposed. Leaving it checked means new tables are public by default until you add RLS — a security footgun.

3. "Enable automatic RLS" — check this. It creates a trigger that automatically enables Row Level Security on new tables. Combined with unchecking the one above, this is the safer default: tables are RLS-protected from creation, not exposed until you explicitly allow it.

4. Region — Americas is fine if your users are primarily US-based, which for V1 they will be.

5. Save that database password somewhere safe right now — hit Copy before you leave this page. Supabase won't show it to you again and you'll need it if you ever connect directly to the database.

6. Skip the GitHub connection for now — Claude Code will handle the repo setup, and you can connect GitHub to Supabase later if you want the auto-deploy feature.

Make those two security checkbox changes, rename the project, then hit Create.

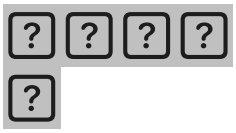

You said: Project Settings → API in supabase is greyed out, i can't click it

Project Settings → API in supabase is greyed out, i can't click it

May 24

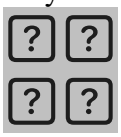

Claude responded: That's normal — Supabase takes 1-2 minutes to provision the project after you create it.

That's normal — Supabase takes 1-2 minutes to provision the project after you create it. The settings are greyed out while it's spinning up the database.

You'll see a progress indicator or a "Setting up your project" message. Just wait for it to finish — when it's done everything will become clickable and you'll land on the project dashboard.

While you're waiting, go ahead and open mux.com in another tab and start that account creation so you're not sitting idle.

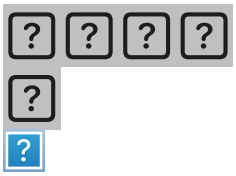

May 24

Claude responded: Project is live.

Project is live. Now click **API Keys** in the left sidebar — that's where your three credentials are.

You'll see:

- **Project URL** → this is your `NEXT_PUBLIC_SUPABASE_URL`
- **anon public** key → `NEXT_PUBLIC_SUPABASE_ANON_KEY`
- **service_role secret** key → `SUPABASE_SERVICE_ROLE_KEY` (click reveal to see it)

Copy all three into a notes app or directly into a `.env.local` file as you go. Don't share them in this chat.

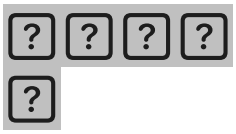

You said: where do i find the Project URL

where do i find the Project URL

May 24

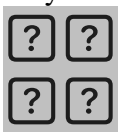

Claude responded: It's on the same API Keys page — scroll up, it'll be at the top of that page above the keys.

It's on the same API Keys page — scroll up, it'll be at the top of that page above the keys. It looks like:

https://iuxhlflulrpoppuwnrxg.supabase.co

Basically `https://[your-project-id].supabase.co` — and you can see your project ID is `iuxhlflulrpoppuwnrxg` right there on the General settings page you're already on. So your URL is already known:

NEXT_PUBLIC_SUPABASE_URL=https://iuxhlflulrpoppuwnrxg.supabase.co

You can fill that one in right now without even navigating. Then click API Keys for the anon and service role keys.

May 24

Claude responded: Supabase has updated their API key naming.

Supabase has updated their API key naming. The mapping to your `.env.local` is:

Publishable key (the `sb_publishable_...` one) → this is your `NEXT_PUBLIC_SUPABASE_ANON_KEY`

Secret key (the `sb_secret_...` one) → this is your `SUPABASE_SERVICE_ROLE_KEY` — click the eye icon to reveal it, then copy it

So your three Supabase env vars are now complete:

```
NEXT_PUBLIC_SUPABASE_URL=https://iuxhlflulrpoppuwnrxg.supabase.co
NEXT_PUBLIC_SUPABASE_ANON_KEY=sb_publishable_xQOfkdLzjHwIcqW-eBzGtA_bel4E...
SUPABASE_SERVICE_ROLE_KEY=sb_secret_womWs...
```

One note for Claude Code — when you hand it these credentials, mention that Supabase has switched to the new `sb_publishable_` / `sb_secret_` key format. Older Supabase client setup guides expect `eyJ...` JWT-format keys. Claude Code may need to know the new format is valid so it doesn't flag them as wrong.

Now head to mux.com and do the same thing there.

May 24

Claude responded: Check these three:

Check these three:

- **Store and stream on-demand videos**
- **Let users upload videos to my application**
- **Use Mux Player to play videos in my app**

Those are exactly what Showside needs — fan uploads, storage, and playback. Skip "live stream" and "monitor performance" for now.

"How did you hear about us" — pick whatever, doesn't matter. Then hit Get Started.

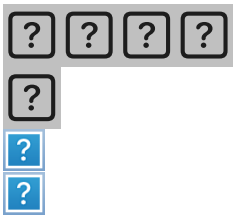

May 24

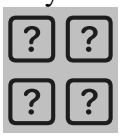

Claude responded: Good — two quick calls:

Good — two quick calls:

Type: Keep it as **Development** for now. You'll create a separate Production environment when you're ready to go live. Development gives you test mode which is free and has no usage limits for building.

Member Access: Keep **All Members** — you're the only member anyway.

Name is already "Showside" which is correct. Hit **Create**.

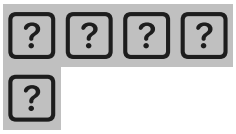

You said: i have a cloudflare account already for my business - should i use that account or create a fresh one

i have a cloudflare account already for my business - should i use that account or create a fresh one

May 24

Claude responded: Use your existing account — no need to create a new one.

Use your existing account — no need to create a new one. R2 is just a service within Cloudflare, completely separate from whatever domains or other services you have set up for RoomBox.

Just go to your existing Cloudflare dashboard, find R2 in the left sidebar, and create a new bucket called `showside-media` there. The billing and API tokens are all scoped to the bucket level so there's no bleed between your business stuff and Showside.

May 24

Claude responded: Go ahead and click "Add R2 subscription to my account" — total due now is $0.

Go ahead and click "**Add R2 subscription to my account**" — total due now is $0.00 and you only get charged if you exceed the free tier (10GB storage, 1M writes, 10M reads per month). You won't get anywhere near those limits in development or early V1.

You need to add the subscription before you can create a bucket, so this is a required step.

May 24

Claude responded: Everything looks correct here:

Everything looks correct here:

- **Bucket name:** `showside-media` ✓
- **Location:** Automatic / Eastern North America ✓ — fine for V1
- **Storage class:** Standard ✓ — photos will be accessed frequently

Hit **Create bucket**.

After it creates, you'll need to do two more things in R2:

1. **Enable public access** on the bucket (so photos are publicly viewable via URL)
2. **Create an API token** with read/write access to get your `R2_ACCESS_KEY_ID` and `R2_SECRET_ACCESS_KEY`

Both are in the bucket settings after creation.

May 24

Claude responded: Bucket is created.

Bucket is created. Two things to do on this page:

1. Enable Public Development URL — click **Enable** next to "Public Development URL". This gives you a public URL to serve photos from during development. It'll generate something like `https://pub-xxxx.r2.dev` — that becomes your `R2_PUBLIC_URL` env var.

2. Note your S3 API URL — you can see it's already showing:

https://eaaff326cceff99d85e30665f5c43281.r2.cloudflarestorage.com/shows...

The account ID portion (`eaaff326cceff99d85e30665f5c43281`) is your `R2_ACCOUNT_ID`. Copy that now.

After you enable the public URL, you need to go back to the main R2 overview page and create an API token. Look for "Manage R2 API Tokens" — it's usually a link on the R2 overview, not inside the bucket settings.

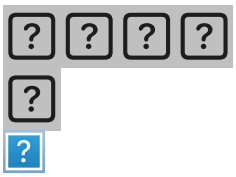

May 24

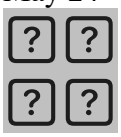

Claude responded: Everything you need is right here on this page.

Everything you need is right here on this page.

Copy your Account ID — click the copy icon next to `eaaff326cceff99d85e30665f5c4...` — that's your `R2_ACCOUNT_ID`.

Create your API token — click **Manage** next to "API Tokens". On the next page, create a new token with:

- Permission: **Object Read & Write**
- Bucket: **Specific bucket → showside-media**

That will give you your `R2_ACCESS_KEY_ID` and `R2_SECRET_ACCESS_KEY`. **Copy both immediately when they're shown — Cloudflare only shows the secret key once.**

Then go back to the bucket settings to enable the Public Development URL if you haven't already — that gives you the `R2_PUBLIC_URL`.

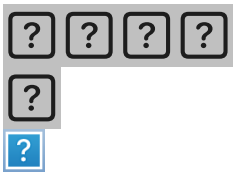

May 24

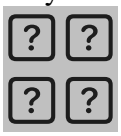

Claude responded: Click "Create Account API token" — the top one, not User API token.

Click "**Create Account API token**" — the top one, not User API token. The description says it's recommended for production systems and stays active even if you leave the org, which is

what you want for an app's server credentials.

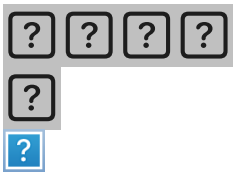

May 24

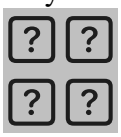

Claude responded: Make these changes before hitting Create:

Make these changes before hitting Create:

Token name: Change from "R2 Account Token" to `showside-media-rw` — just so it's clear what it's for when you see it later.

Permissions: Select **Object Read & Write** (third option) — not the currently selected "Object Read only". You need write access to upload photos.

Specify buckets: Switch to **"Apply to specific buckets only"** and select `showside-media`. Don't give it access to all buckets — least privilege, especially since this is your business Cloudflare account.

TTL: Forever is fine.

Client IP filtering: Leave blank.

Then hit **Create Account API Token** and copy both keys immediately when shown.

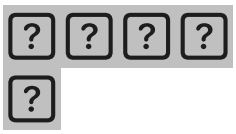

You said: now what

now what

May 24

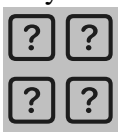

Claude responded: Token is created and active.

Token is created and active. Now you need to go back and enable the Public Development URL on the bucket so photos are publicly accessible.

Go to: **Storage & databases → R2 Object Storage → showside-media → Settings** and click **Enable** next to "Public Development URL". It'll generate a URL like `https://pub-xxxxxxxxxxxx.r2.dev` — that's your `R2_PUBLIC_URL`.

Once you have that, your complete `.env.local` should look like this — fill in the blanks with what you've collected:

bash

```
# Supabase
NEXT_PUBLIC_SUPABASE_URL=https://iuxhlflulrpoppuwnrxg.supabase.co
NEXT_PUBLIC_SUPABASE_ANON_KEY=sb_publishable_...
SUPABASE_SERVICE_ROLE_KEY=sb_secret_...

# Mux
MUX_TOKEN_ID=
MUX_TOKEN_SECRET=

# Cloudflare R2
R2_ACCOUNT_ID=eaaff326cceff99d85e30665f5c43281
R2_ACCESS_KEY_ID=        ← from the token you just created
R2_SECRET_ACCESS_KEY=    ← from the token you just created (only shown once)
R2_BUCKET_NAME=showside-media
R2_PUBLIC_URL=           ← from the public dev URL you're about to enable

# Admin
ADMIN_KEY=               ← make up any long random string, e.g. "showside-admin-2026"

# App
NEXT_PUBLIC_APP_URL=http://localhost:3000
```

After the public URL, the only thing left is Mux — go grab those two keys from the Mux dashboard and you're done with credentials.

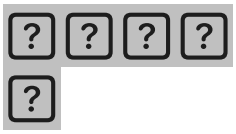

You said: where do i find the

NEXTPUBLICSUPABASE_URL

where do i find the NEXT_PUBLIC_SUPABASE_URL

May 24

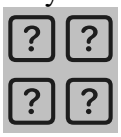

Claude responded: You already have it — it's just your Supabase project ID formatted as a URL:

You already have it — it's just your Supabase project ID formatted as a URL:

https://iuxhlflulrpoppuwnrxg.supabase.co

You can see the project ID `iuxhlflulrpoppuwnrxg` right in the Supabase dashboard URL in your earlier screenshot: `supabase.com/dashboard/project/iuxhlflulrpoppuwnrxg/...`

So that env var is already filled in — no need to look it up.

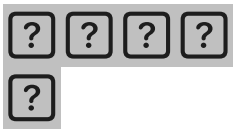

You said: ok i think i have everything - in a text document - paste it into claude code?

ok i think i have everything - in a text document - paste it into claude code?

May 24

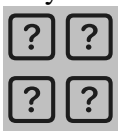

Claude responded: Yes — paste it directly into Claude Code with this message:

Yes — paste it directly into Claude Code with this message:

Here are my credentials. Please create the `.env.local` file with these values and proceed to Phase 1.

[paste your env vars here]
Note: Supabase is using their new API key format — the publishable key starts with `sb_publishable_` and the secret key starts with `sb_secret_`. These are valid, not the old `eyJ...` JWT format. The Supabase JS client v2 accepts both formats.

That context note at the end is important — Claude Code may have been trained on docs showing the old JWT-format keys and might flag yours as wrong without it.

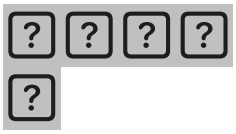

You said: i started it running opus 4.

i started it running opus 4.7 - should i stop it and run sonnet instead or just let it run?

May 24

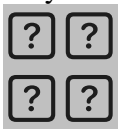

Claude responded: Let it run.

Let it run. You're in exactly the right situation for Opus — Phase 1 is the architecture, schema, and foundation decisions. That's precisely where the extra reasoning pays for itself. Opus thinking through your RLS policies, trigger logic, and project structure correctly now saves you from untangling bad decisions later.

Switch to Sonnet when you hit the repetitive implementation work — building out the individual pages and components in Phases 4 and 5. That's when Opus stops being worth the extra cost and Sonnet handles it just as well.

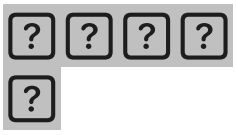

You said: phase 1 response

I'm in the Google Drive spec folder. I'll create the Next.js project in a web/ subfolder so the

spec docs stay alongside as reference. One quick note before I dive in: Google Drive can corrupt node_modules and fight with file watchers during active dev. I'll proceed since this is where you're set up

pasted

phase 1 response

May 24

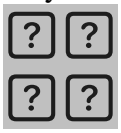

Claude responded: Really solid Phase 1.

Really solid Phase 1. A few things to action right now:

1. Move the project out of Google Drive first — do this before anything else.

This is the most important thing in the whole response. Your `.env.local` with real Supabase, Mux, and R2 credentials is currently syncing to Google's servers. Open Terminal and run:

bash
```
mv ~/Google\ Drive/My\ Drive/[your-spec-folder]/web ~/Projects/showside
```
Create `~/Projects/` first if it doesn't exist:

bash
```
mkdir -p ~/Projects
mv ~/path/to/gdrive/web ~/Projects/showside
```
Then reopen Claude Code pointed at the new location.

2. Apply the migration and seed to Supabase.

Go to: `https://supabase.com/dashboard/project/iuxhlflulrpoppuwnrxg/sql/new`

Open the file `web/supabase/migrations/0001_init.sql`, copy the entire contents, paste into the SQL editor, run it. Then do the same for `web/supabase/`

`seed.sql`.

3. The open questions from last time are already answered.

Claude Code says items 3–8 are still open but you answered all of them earlier in our conversation. When you tell it to proceed to Phase 2, paste this reminder:

All design questions were already answered: (1) reword "close this tab" to "leave this tab open"; (2) magic link signin at /signin not in nav; (3) Featured with cumulative view_count; (4) like toggle = unlike on repeat; (5) store aspect_ratio from Mux metadata; (6) orphan rows use status 'uploading' as you've already implemented. Proceed to Phase 2.

Once the migration is applied and the project is moved, tell it to start Phase 2.

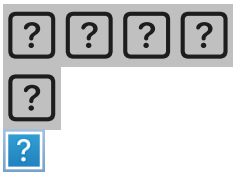

May 24

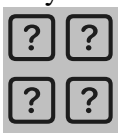

Claude responded: The path I gave you used a placeholder — you need to find the actual path.

The path I gave you used a placeholder — you need to find the actual path. Run this to see what's in your Google Drive:

bash

```
ls ~/Library/CloudStorage/
```

On modern Macs Google Drive lives there rather than `~/Google Drive`. You'll see something like `GoogleDrive-mpocock@orderroombox.com`. Then run:

bash

```
ls ~/Library/CloudStorage/GoogleDrive-mpocock@orderroombox.com/My\ Drive/
```

That'll show you the actual folder name where the project lives. Once you see it, the move command will be:

bash
```
mkdir -p ~/Projects
mv ~/Library/CloudStorage/GoogleDrive-mpocock@orderroombox.com/My\ Drive/[actual-folder-name]/web ~/Projects/showside
```
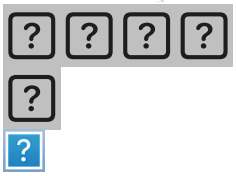
May 24
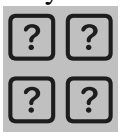

Claude responded: I can see your Google Drive root but the project folder isn't visible here — it's probably named something like "Showside" or "showside-spec".

I can see your Google Drive root but the project folder isn't visible here — it's probably named something like "Showside" or "showside-spec". Scroll up in that listing or run this to find it:

bash
```
ls ~/Library/CloudStorage/GoogleDrive-mpocock@orderroombox.com/My\ Drive/ | grep -i show
```
Or if Claude Code created it with a different name:

bash
```
ls ~/Library/CloudStorage/GoogleDrive-mpocock@orderroombox.com/My\ Drive/ | grep -i web
```
That'll narrow it down. Once you see the folder name, paste it here and I'll give you the exact move command.

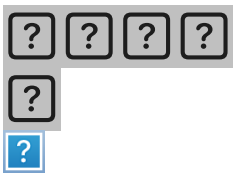
May 24
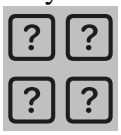

Claude responded: The grep returned nothing, which means Claude Code created the project somewhere other than your Google Drive root.

The grep returned nothing, which means Claude Code created the project somewhere other than your Google Drive root. It might be in a subfolder. Run this to find it:

bash
```
find ~/Library/CloudStorage/GoogleDrive-mpocock@orderroombox.com -name "package.json" -path "*/web/*" 2>/dev/null
```
That'll locate the exact path of the project wherever it ended up.

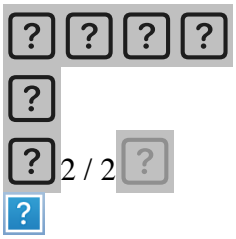

Claude responded: The find returned nothing — the project isn't in Google Drive at all.

The find returned nothing — the project isn't in Google Drive at all. Claude Code created it somewhere else. Run this to find it:

bash
```
find ~ -name "package.json" -path "*/showside/*" 2>/dev/null
```
Or check the Claude Code project context — look at the bottom of the Claude Code window where it shows the project path. It'll show you exactly where the `web/` folder lives.

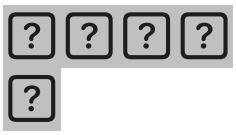

You said: /Users/mpocock/Library/CloudStorage/

GoogleDrive-mpocock@orderroombox.

/Users/mpocock/Library/CloudStorage/GoogleDrive-mpocock@orderroombox.com/My Drive/Personal G Drive/Biz Ideas/Showside

May 24

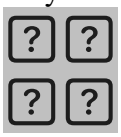

Claude responded: Found it.

Found it. Run this exact command:

bash
```
mv "/Users/mpocock/Library/CloudStorage/GoogleDrive-mpocock@orderroombox.com/My Drive/Personal G Drive/Biz Ideas/Showside/web" ~/Projects/showside
```
Then verify it moved correctly:

bash
```
ls ~/Projects/showside
```
You should see the project files — `package.json`, `app/`, `lib/`, etc.

After that, reopen Claude Code and point it at `~/Projects/showside` as the project directory before continuing to Phase 2.

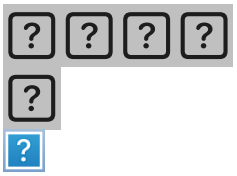

May 24

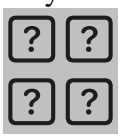

Claude responded: Phase 1 is fully verified and it's ready for Phase 2.

Phase 1 is fully verified and it's ready for Phase 2. Everything looks great — schema is in

Supabase, seed data is there, build is clean.

Two things before you hit go:

1. Confirm the like toggle — it's asking one more time. The answer is yes, repeat call = unlike. Tell it that.

2. Did you move the project out of Google Drive? The screenshot shows it's still in the Showside project context but I can't tell if the move happened. Check the project path at the bottom of Claude Code — if it still shows the Google Drive path, do the move first before proceeding.

If the move is done, reply with:

Project is moved to ~/Projects/showside. Like toggle confirmed: repeat call = unlike. Go ahead with Phase 2.

If you haven't moved it yet, do that first, repoint Claude Code at the new location, then send that message.

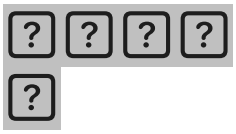

You said: where am i moving it to - my local device?

where am i moving it to - my local device?

May 24

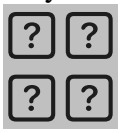

Claude responded: Yes — just to a regular folder on your Mac's local hard drive, not inside Google Drive.

Yes — just to a regular folder on your Mac's local hard drive, not inside Google Drive. `~/Projects/showside` means:

/Users/mpocock/Projects/showside

That folder lives only on your MacBook, not synced anywhere. It's the standard place developers keep code on a Mac.

The reason is purely practical — Google Drive actively fights with Next.js development. It tries

to sync `node_modules` (500+ files changing constantly), conflicts with the file watcher that powers hot reload, and your `.env.local` with real credentials gets uploaded to Google's servers.

The spec docs you want to keep as reference can stay in Google Drive. Just the `web/` code folder needs to move to local.

Run this if you haven't already:

bash
```
mv "/Users/mpocock/Library/CloudStorage/GoogleDrive-mpocock@orderroombox.com/My Drive/Personal G Drive/Biz Ideas/Showside/web" ~/Projects/showside
```
Then reopen Claude Code pointed at `~/Projects/showside`.

You said: ok i created a folder on my desktop - how do i do the move

ok i created a folder on my desktop - how do i do the move

May 24

Claude responded: I can see you created a "Showside App" folder on your Desktop.

I can see you created a "Showside App" folder on your Desktop. Run this in Terminal:

bash
```
mv "/Users/mpocock/Library/CloudStorage/GoogleDrive-mpocock@orderroombox.com/My Drive/Personal G Drive/Biz Ideas/Showside/web" "/Users/mpocock/Desktop/Showside App/web"
```
Then verify:

bash
```
ls "/Users/mpocock/Desktop/Showside App/web"
```
You should see `package.json`, `app/`, `lib/`, etc.

After that, go back to Claude Code, click the project selector at the top, and point it at `/Users/mpocock/Desktop/Showside App/web`.

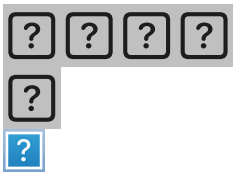

May 24

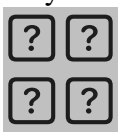

Claude responded: The commands ran without errors which is good, but I can't see the output of the final ls — scroll down in Terminal to see if it listed the project files.

The commands ran without errors which is good, but I can't see the output of the final `ls` — scroll down in Terminal to see if it listed the project files.

Run this to confirm:

bash
```
ls "/Users/mpocock/Desktop/Showside App/web"
```
If you see `package.json`, `app`, `lib`, `node_modules` etc. — the move worked and you're good to repoint Claude Code.

If it says "No such file or directory" — the project was never in Google Drive to begin with and Claude Code created it somewhere else. In that case run:

bash
```
find ~ -name "next.config.js" 2>/dev/null | grep -v node_modules
```
That'll find it wherever it actually is.

You said: when i run that first command nothing comes up.

when i run that first command nothing comes up...

May 24

Claude responded: Run this to find where it actually is:

Run this to find where it actually is:

bash

```
find ~ -name "next.config.js" 2>/dev/null | grep -v node_modules
```

That'll locate it wherever Claude Code created it.

You said: a bunch of access prompts started popping up

a bunch of access prompts started popping up

May 24

Claude responded: It already moved — look at the bottom of the terminal output:

It already moved — look at the bottom of the terminal output:

```
/Users/mpocock/Projects/showside/next.config.js
```

And earlier you can see `ls ~/Projects/showside` returned all the project files:

app middleware.ts node_modules package.json lib next.config.js etc...

The project is at `~/Projects/showside`. The Desktop move failed but the first `~/Projects/showside` move worked perfectly.

Dismiss all those access permission popups — they were triggered by the find command scanning your whole home directory. Just click "Don't Allow" or close them.

Then go to Claude Code, click the project selector at the top, and point it at:

/Users/mpocock/Projects/showside

You said: where exactly do i click ?

where exactly do i click ?

May 24

Claude responded: You need to change the working directory Claude Code is pointed at.

You need to change the working directory Claude Code is pointed at. Type this directly in the Claude Code chat:

The project has been moved to ~/Projects/showside. Please use that as your working directory going forward. Like toggle confirmed: repeat call = unlike. Go ahead with Phase 2.

Claude Code will pick up the new path from that message and work from the correct location. You don't need to click anything — just tell it where the project is.


May 24


Claude responded: Perfect — it picked it up immediately.

Perfect — it picked it up immediately. You can see:

"Project verified at ~/Projects/showside/ with node_modules intact. Creating Phase 2 tasks"

It's now working from the correct local directory, out of Google Drive. Phase 2 is underway. Just let it run — it'll build all 8 API routes and test them with curl before any UI work. This is a good time to step away and let Opus do its thing.

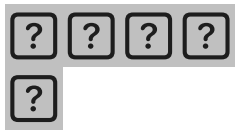

You said: Confirmed.

Confirmed. First verifying the project landed correctly at the new path, then loading task tools (Phase 2 has 8 routes — worth tracking), then starting: Ran a command, used a tool Project verified at ~/Projects/showside/ with node_modules intact. Creating Phase 2 tasks: Created 8 files, used 23 tools All 8 routes written. Marking #8 done and running typecheck + build to verify: Ran a command, used a tool Two type errors to fix. The exit code reading is misleading (it's head's exit, not tsc's — pipeline issue). Fixing: Edited 2 files, ran 2 commands Typecheck clean. Build smoke test: Ran Build with new routes All 8 routes registered. Now smoke-testing with curl. Starting the dev server in background: Ran 2 commands Running first batch of tests in parallel — search, QR, admin login (wrong + right): Ran 3 commands Admin login works. Search and QR both 500'd — need to check the dev server logs to see what's happening: Ran Read dev server logs Need to see the actual error. Testing Supabase connectivity directly + checking what's installed: Ran 2 commands Found it — the service_role doesn't have table grants. Standard Supabase auto-grants via the dashboard, but tables created through the SQL editor don't get auto-granted. Writing a fix migration: Created a file, ran a command The grants migration (34 lines) is on your clipboard. Same drill in the Supabase SQL editor: New query (+ button) or clear current Paste (Cmd+V) Run Expected: "Success. No rows returned." Once that's done, tell me and I'll re-run the curl tests — search and QR should both work. Admin login was already passing. While you do that, here's where Phase 2 stands: Endpoint Status POST /api/admin/login ✅ Verified (401 + 200 + cookie) GET /api/search ⏸ Blocked on grants GET /api/events/[slug]/qr ⏸ Blocked on grants POST /api/upload/photo-url ⏸ Pending test (needs grants) POST /api/upload/video-url ⏸ Pending test (needs grants) POST /api/upload/complete ⏸ Pending test (needs grants) POST /api/media/[id]/like ⏸ Pending test (needs grants) POST /api/mux/

webhook ⏸ Will document — needs ngrok + signed payload to test All routes typecheck and build cleanly; this is a database-side permissions issue, not a code issue.

Show more
May 24

Claude responded: Good progress.

Good progress. Go to the Supabase SQL editor:

`https://supabase.com/dashboard/project/iuxhlflulrpoppuwnrxg/sql/new`

Click the + button for a new query, paste the migration from your clipboard, hit Run. You should see "Success. No rows returned."

Then come back and tell Claude Code:

Grants migration applied. Re-run the curl tests.

You said: Grants applied.

Grants applied. Re-running the full smoke test suite: Ran Retest search and QR after grants Both green. Now the upload → complete → like flow: Ran End-to-end upload flow test Phase 2 complete and verified end-to-end. All seven testable routes pass: Test Result GET /api/search?q=morgan ✅ Returns Morgan Wallen entity + 3 events with embedded entity object GET /api/events/[slug]/qr ✅ Returns valid 512×512 PNG (3.9 KB) POST /api/admin/login (wrong) ✅ 401 POST /api/admin/login (right) ✅ 200 + showside_admin cookie set POST /api/upload/photo-url ✅ Returns mediaId + R2 presigned PUT URL (5 min) + canonical storage URL POST /api/upload/complete ✅ Tags applied, status flipped uploading → active POST /api/media/[id]/like (1st) ✅ { liked: true, likeCount: 1 } POST /api/media/[id]/like (2nd) ✅ { liked: false, likeCount: 0 } — toggle confirmed POST /api/upload/video-url ✅ Returns Mux direct-upload URL + muxUploadId The DB triggers fired correctly too — like_count went 0 → 1 → 0 cleanly, which means the trigger on the likes table is working and using the same transaction as the insert/delete. Mux webhook deferred from smoke testing — it needs: A public URL Mux can

POST to (ngrok or Vercel preview) A signed payload (only Mux can generate one with our secret) The code is in place and verifies HMAC-SHA256 with replay protection. To smoke-test it, you'd: install ngrok → ngrok http 3000 → set the public URL in your Mux dashboard webhook config → upload a real video via the video-url endpoint → watch the webhook fire. I can walk you through that in detail when we're ready — happy to do it now, or punt until we have UI to upload from. One cosmetic issue spotted — the seed event names show up with mojibake (e.g. "Morgan Wallen ,Äî AT&T Stadium"). That's the em-dash — from my seed.sql getting mis-decoded between clipboard and Supabase. Not blocking, easy to fix with an UPDATE later. Flagging so it doesn't surprise you on the entity pages. Dev server still running in the background (task ID bf33hysgl). You can keep using it. To stop it I'd need to kill it explicitly — say the word and I will. Ready for Phase 3 Phase 3 is the admin panel — needed before public pages so you can create entities/events through a UI instead of pasting SQL. Scope: / admin dashboard (stats + nav) /admin/entities (list + create + edit) /admin/events (list + create + edit + paste-setlist textarea) /admin/moderation (queue of pending_review media + approve/ remove) QR code download button per event Want me to: Fix the em-dash mojibake first (one quick UPDATE), then start Phase 3, or Start Phase 3 directly and clean up the mojibake whenever, or Test the Mux webhook with ngrok now before moving on?

Show more
May 24

Claude responded: Great progress — all 7 testable routes green, triggers working, Mux upload URL confirmed.

Great progress — all 7 testable routes green, triggers working, Mux upload URL confirmed.

Tell it:

Fix the em-dash mojibake first, then go straight into Phase 3. Skip the Mux webhook ngrok test for now — we'll verify it when we have real UI to upload from.

You said: when can i see a mockup of it

when can i see a mockup of it

May 24

Claude responded: You can see it running locally right now actually — the dev server is already running in the background.

You can see it running locally right now actually — the dev server is already running in the background.

Open your browser and go to:

http://localhost:3000
It'll just be the Phase 1 placeholder page at this point — nothing designed yet. The real UI starts coming together in Phase 4 (read pages) after the admin panel is done in Phase 3.

So the timeline is roughly:

- **Phase 3 (now):** Admin panel — functional but unstyled, just for you to create entities/events
- **Phase 4 (next):** The public-facing pages start getting built — entity page, event page, home, viewer. This is when it starts looking like the mockups we designed
- **Phase 5:** Upload flow UI
- **Phase 6:** SEO polish and final Lighthouse pass

So realistically after Phase 4 is done you'll have something that looks like an actual product. Based on how fast Phase 2 went, that's probably another 2-3 Claude Code sessions from now.

You said: When you're ready to move on, I still need answers on the three Phase 4 questions before I start:

When you're ready to move on, I still need answers on the three Phase 4 questions before I start:

1. Returning-user sign-in — ship `/signin` (magic link, not in nav) or skip entirely?
2. "Most-viewed this week" hero — cumulative `view_count` relabeled "Featured", or add `weekly_views`

    - cron?
3. "Stage moments" wide-angle — drop the criterion, or store Mux aspect-ratio metadata?

Or just say "go with your defaults" and I'll proceed with: no `/signin`, "Featured" cumulative, no aspect-ratio storage.

May 24

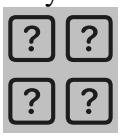

Claude responded: This is looking great — admin panel is fully functional, all 4 entities and 9 events are showing up, the em-dash issue is visible but minor.

This is looking great — admin panel is fully functional, all 4 entities and 9 events are showing up, the em-dash issue is visible but minor.

Tell it:

Go with your defaults on all three. No /signin for now, Featured with cumulative view_count, no aspect ratio storage. Fix the em-dash mojibake then start Phase 4.

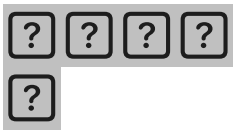

You said: will i get to see what the frontend looks like for users also?

will i get to see what the frontend looks like for users also? how is the design/UI being determined?

May 24

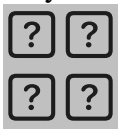

Claude responded: Yes — Phase 4 is exactly that.

Yes — Phase 4 is exactly that. The public-facing pages (entity page, event page, home, viewer) get built next and you'll be able to see them at localhost:3000.

For the design — right now Claude Code is making its own decisions on the UI since the scope document described functionality but didn't specify a visual design system. It'll produce

something functional but probably generic.

You have two options:

Option A — Let it build, then redesign. Let Claude Code finish the functional build first, then in a separate pass give it the mockups we designed in this conversation and tell it to restyle everything to match. This is the faster path to a working product.

Option B — Feed it the design direction now. Before it starts Phase 4, paste it the mockup specs from our conversation — the dark cinematic hero, the experience-first entity page layout, the highlights grid, the color palette. It'll build closer to what we designed from the start. More upfront work but less rework later.

I'd recommend **Option B** given how much design thinking we've already done. When you tell it to start Phase 4, add something like:

Before building the public pages, here is the design direction: dark theme, cinematic feel, experience-first not archive-first. The entity page should lead with a dark hero section containing video thumbnails on the right and artist name/stats on the left. Below that: fan highlights grid (1 large + 4 small), most filmed songs row, recent shows row, then a collapsed full archive. Reference the attached mockup descriptions for layout details.

Want me to write you a proper design brief to paste in alongside that?

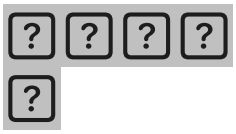

You said: lets see what it comes back with.

lets see what it comes back with. i want to research some designs and then we can update it later

May 24

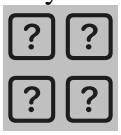

Claude responded: Smart approach.

Smart approach. Let Claude Code build the functional foundation first, then you can come back with specific design direction once you've done your research and know exactly what you want.

Tell it:

Go with your defaults on all three. Fix the em-dash mojibake then start Phase 4. Use your best judgment on design for now — we'll do a full design pass later.

Then let it run. Come back when Phase 4 is done and we can look at what it produced together and plan the redesign from there.

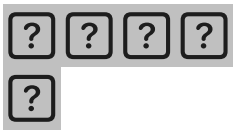

You said: Building all remaining Phase 4 files in one batch — home, search, entity, SEO:

Building all remaining Phase 4 files in one batch — home, search, entity, SEO: Created 6 files, read a file, used 6 tools Adding JSON-LD to the event page and canonical URL to metadata: Edited 2 files, ran a command, used 2 tools Smoke-testing all public pages (skipping next build this time to avoid stomping the dev server — verifying via the live dev server instead): Ran Smoke-test all Phase 4 public pages Quick sanity check on the artist filter — I think Kentucky Derby is appearing only in the "Most uploaded shows" section, not the filtered entity grid: Ran a command, used a tool Phase 4 complete. All 8 tasks (15–22) done. The full public read surface is shipping clean. Smoke-test results: Page Result / (home) ✅ Renders Trending + Most uploaded shows + Browse by type, all 4 entities visible /?type=artist ✅ Filters entity grid to Morgan Wallen + Taylor Swift only /?type=team ✅ Only Phillies /?type=event_brand ✅ Only Kentucky Derby /search ✅ Loads with empty state /search?q=swift ✅ Loads (client component will debounce-fetch on render) /morgan-wallen ✅ Hero, Fan highlights (5 filter pills), Most filmed songs, Recent shows, Full archive /morgan-wallen?filter=recent ✅ Filter applied via URL /morgan-wallen/morgan-wallen-att-stadium-dallas-2026-07-18 ✅ Watch/Browse/Upload tabs /watch/<active mediaId> ✅ Full-screen viewer /sitemap.xml ✅ Includes all 4 entities + 9 events /robots.txt ✅ Disallows /admin/, /api/, /upload/, points to sitemap JSON-LD on /morgan-wallen ✅ MusicGroup structured data emitted JSON-LD on event page ✅ Event schema with location + performer Things you can try in the browser right now: http://localhost:3000 — discovery feed http://localhost:3000/morgan-wallen — most complex page; click filter pills, scroll down to Full archive and expand it http://localhost:3000/search — type "swift" or "phillies" http://localhost:3000/sitemap.xml — see what Google would crawl Did NOT run next build — that's the gotcha that broke us earlier (build stomps dev server's .next/). The dev server's own HMR + tsc --noEmit cover the same ground. I'll do a clean next build only when the dev server is intentionally stopped. What's left in V1 Phase 5 — Upload flow (the last unbuilt user-facing thing): /upload — event picker with autocomplete against the events table /upload/[event-slug] — single-page stepper (event confirm → file picker → optional tagging → direct upload to Mux/R2 → success screen with optional account claim) This is the most UX-critical flow per spec. Mobile-first. Has to feel fast. Will give us a way to finally end-to-end test the Mux webhook with a real video. Polish / V1 acceptance items not yet covered: View-count dedup (currently increments every page load — flagged from the start) Image optimization (we're using unoptimized on Next/Image — flip when R2 + Mux remote patterns are confirmed

in next.config.js, which they already are) Virtualization on long media grids (spec says virtualize > 50 items — current pages cap at 100/200, fine for V1) Infinite scroll (currently fixed limit; acceptance criteria #2 doesn't require it) Lighthouse > 80 mobile audit Footer pages (/about, /dmca, etc. — currently 404) Want me to start Phase 5 (the upload flow), or pause here so you can poke around the site and decide?

Show more
May 25

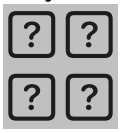

Claude responded: Go to your browser and check it out first:

Go to your browser and check it out first:

- `http://localhost:3000` — home page
- `http://localhost:3000/morgan-wallen` — the entity page, click the filter pills and expand the full archive
- `http://localhost:3000/search` — try typing "swift" or "phillies"

Take a few minutes to click around and see what it built. Then come back and let me know what you think — what looks good, what's missing, what feels off. We can decide together whether to push straight into Phase 5 or do the design pass first before adding more functionality.

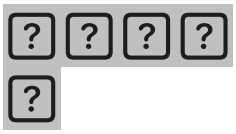

You said: write a summary of this chat i can use to feed a new chat t help me preserve tokens

write a summary of this chat i can use to feed a new chat t help me preserve tokens

May 25

Claude responded: Here's a summary you can paste at the start of a new chat:

Here's a summary you can paste at the start of a new chat:

Showside — Project Context

What we're building

Showside — a fan media platform for live events. Fans who attended upload photos/videos. Non-attendees virtually attend by browsing organized fan footage. Think YouTube Shorts meets a permanent concert archive.

Core concept: Experience-first, archive-second. A teenager in London who couldn't afford a Morgan Wallen ticket lands on his page and feels like she experienced the show within 30 seconds. The entity page leads with the best fan clips across all shows — not a list of tour dates.

Tagline: "The show, from everyone who was there."

Product decisions made

- **Primary audience:** People who couldn't attend (priced out, wrong city, etc.) — not just attendees
- **Entity → Event hierarchy:** Artist/team/brand page → specific show sub-page → media grid
- **Entity page design:** Experience-first. Hero with cinematic dark background + video thumbnails. Fan highlights grid (1 large + 4 small cards). Most filmed songs row. Recent shows (4 cards). Full archive collapsed behind a click.
- **Video dominates** over photos — virtual attendance needs video
- **Setlist-to-clip linking:** Manual tag by uploader + AI audio fingerprinting (ACRCloud) in V2
- **Anonymous uploads** — no account required to upload. Account creation only prompted on success screen.
- **Upload target:** Under 60 seconds from opening page to upload complete on mobile
- **Like toggle:** Repeat call = unlike
- **No audio fingerprinting in V1** — manual song tagging only
- **No mobile app in V1** — mobile-first web only
- **Mux for video** (upload, transcode, HLS, thumbnails), **Cloudflare R2 for photos**
- **Featured** (cumulative view_count) not "most viewed this week" — no cron job needed
- **No /signin in V1** — magic link deferred
- **No aspect ratio storage in V1**
- **Orphan media rows:** status = 'uploading' on presign, flipped to 'active' on complete/webhook

Tech stack

- **Framework:** Next.js 14 (App Router, SSR)
- **Database:** Supabase (Postgres + auth + RLS)

- **Video:** Mux (direct upload, HLS playback, webhooks)
- **Photos:** Cloudflare R2 (presigned URLs, public dev URL)
- **Deployment:** Vercel (not yet deployed)
- **Styling:** Tailwind CSS
- **Language:** TypeScript

Competitive landscape

- **CrowdAlbum** — acquired by Spotify 2016, shut down. Aggregated social posts, never tested native direct upload.
- **Concert Archives** — closest existing product. Live since 2013, got NYT/Pitchfork press, but looks dated, sparse media despite large event database, built for attendees logging history not non-attendees experiencing shows.
- **FanVids / Aftershow.io / SureShot / Happaning** — all B2B/artist-controlled. Fans upload to artist's marketing CRM, not a community archive.
- **The gap:** Consumer-first, community-owned, virtual attendance experience. Nobody owns it.

Build status

Using **Claude Code** with **Opus 4.7 + Max effort**. Project located at `~/Projects/showside`.

Completed phases:

- ✅ Phase 1 — Next.js scaffold, Supabase schema + migrations + RLS + triggers, seed data, all lib clients (Supabase, Mux, R2, session), middleware
- ✅ Phase 2 — All 8 API routes built and curl-tested: photo-url, video-url, upload-complete, mux-webhook, like-toggle, search, QR, admin-login
- ✅ Phase 3 — Admin panel: dashboard (stats), entities CRUD, events CRUD + setlist paste, moderation queue, QR download
- ✅ Phase 4 — All public pages: home, entity page, event page (Watch/Browse/Upload tabs), full-screen viewer, search, sitemap.xml, robots.txt, JSON-LD structured data

Remaining:

- ⏳ Phase 5 — Upload flow UI (event picker → file select → optional tagging → direct upload → success screen)
- ⏳ Phase 6 — SEO polish, view-count dedup, infinite scroll, Lighthouse > 80, footer pages (/about, /dmca, etc.)
- ⏳ Design pass — Claude Code used its own judgment on UI. A full design redesign is planned after functionality is complete.

Database schema (summary)

Six tables: `entities`, `events`, `media`, `users`, `likes`, `follows`. DB triggers maintain all counters (upload_count, like_count, follower_count) in-transaction. Anon likes use partial unique index on (media_id, session_token) where user_id is null. Media status flow: `uploading` → `active` (photos on /complete, videos on Mux webhook).

Supabase project

- Project ID: `iuxhlflulrpoppuwnrxg`
- Region: us-west-2

Key product insights from research

- 56% of concert attendees share media on social during shows (Nielsen)
- Concert ticket prices up 50% since 2019 — avg top-100 tour ticket $136 in 2024
- 151M people attended Live Nation events in 2024
- Fan-shot concert clips on YouTube regularly get 10M+ views — demand is proven, just unorganized
- Upload behavior proven by B2B platforms (SureShot, FanVids) when organizer promotes it
- The "virtual attendance session" (30+ min browsing a full show) is the one unproven hypothesis V1 needs to test

Design direction (for future design pass)

- Dark, cinematic aesthetic
- Experience-first not archive-first on entity pages
- Video is the hero content type
- Filter pills on entity page: All time best / Crowd energy / Stage moments / Acoustic / Most recent
- Archive is real and accessible but behind a click — not the front door
- Mobile-first upload experience is the make-or-break UX moment

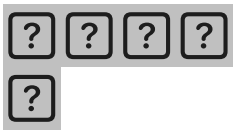

You said: write a detailed description of the showside platform as you understand it today - i want to use this as the Instructions for this Project

write a detailed description of the showside platform as you understand it today - i want to use

this as the Instructions for this Project

May 25

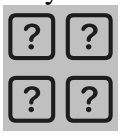

Claude responded: Showside — Project Instructions

Showside — Project Instructions

What Showside is

Showside is a fan media platform for live events. People who attend concerts, sports games, races, and other live events upload their photos and videos to dedicated event pages. Everyone else — people who couldn't get tickets, couldn't afford them, live in another city or country, or simply want to relive the experience — watches, browses, and virtually attends through the fan-captured footage.

The platform sits at the intersection of two proven behaviors: fans already film everything at live events and post it to social media, and non-attendees already watch that content in huge numbers on YouTube and TikTok. What doesn't exist is an organized, permanent, structured destination for that content. Showside is that destination.

The one-line pitch: **"The show, from everyone who was there."**

The problem it solves

Concert ticket prices have risen 50% since 2019. The average top-100 tour ticket hit $136 in 2024. Taylor Swift resale tickets averaged over $1,000. One in five Gen Z fans spent beyond their means just to attend a show. For every person inside a venue, many more wanted to be there and couldn't.

When those fans search YouTube for "Taylor Swift Philadelphia concert," they get a random pile of clips uploaded by different people, in no particular order, with no structure — no way to find every clip from that specific night, no setlist context, no ability to browse by seat section or song. The content exists. The organization doesn't.

Showside takes that same raw material — fan-shot phone footage — and turns it into a structured, navigable show experience for the people who weren't there.

The core mental model

Experience-first, archive-second.

The primary user is someone who didn't attend. She's a 16-year-old in London who loves Morgan Wallen and will never afford a stadium ticket. She lands on his Showside page and within 30 seconds she's watching fan-shot video from the floor pit of last week's Boston show. She doesn't need to pick a specific date or venue to get there. The best content surfaces automatically.

The archive — every specific show, every setlist, every upload organized by date and venue — is real, valuable, and fully accessible. It's just not the front door. It lives behind a single click for the fans who want that level of detail: the superfan who knows exactly which night Morgan played a surprise cover, the attendee who wants to find their own uploads, the person who wants to compare night 1 vs night 2 of a two-night stand.

This is the key product distinction that separates Showside from everything that has been tried before. Concert Archives (the closest existing competitor, launched 2013) is a diary — built for attendees to log their own history. Showside is a broadcast — built for everyone who wasn't there.

How it works

The three-tier hierarchy

Tier 1 — Entity pages (`/morgan-wallen`, `/philadelphia-phillies`, `/kentucky-derby`) Permanent, evergreen pages for any recurring subject that generates live events. Types include: music artists, sports teams, recurring event brands (Kentucky Derby, Indy 500, Coachella), and venues. The entity page is the experience layer — it surfaces the best fan content across all shows, all time, without requiring visitors to pick a specific date first.

Tier 2 — Event pages (`/morgan-wallen/morgan-wallen-gillette-stadium-boston-2026-05-10`) Each specific show, game, or event gets a permanent URL. Tied to a venue, date, city, and optionally a tour name and setlist. This is where uploads attach. The event page has three views: Watch (curated experience for non-attendees), Browse (full unfiltered grid), and Upload (the fan contribution flow).

Tier 3 — Media Individual photos and videos uploaded by fans. Each piece of media can be tagged to a specific song from the event's setlist, a seat section (floor, section 100s, upper deck, stage left/right), and an optional caption. Media items have like counts, view counts, and uploader attribution (anonymous or named).

The entity page in detail

The entity page is the most important page on the platform. It must deliver the virtual attendance experience without requiring any navigation decisions from a casual visitor.

Hero section: Dark, cinematic background using a real fan photo or gradient. Right side shows a 2×3 grid of the top 6 most-viewed video thumbnails for this entity — real content, immediately visible, before any scroll. Left side overlaid on the dark background: entity name, genre badge, verified badge if applicable, follow button, follower count. Below the hero: a stat strip showing total photos & videos / shows covered / contributors / earliest show year.

Fan highlights section: The core of the page. A filter pill row lets visitors sort by: All time best / Crowd energy / Stage moments / Acoustic / Most recent. Below the pills: one large hero card (spans two rows) plus four smaller cards in a 2×2 grid. Each card shows a video thumbnail, play button, duration, song tag or caption, like count, uploader handle, and venue/city. This section answers "what is a Morgan Wallen concert like?" without requiring any further navigation.

Most filmed songs section: A 5-column grid showing the songs that have the most fan clips across all shows — song name, clip count, a relative bar showing volume. Clicking a song filters the page to show all clips of that song from any show. This is a crucial feature: a fan who wants to see every clip of "Last Night" from any show anywhere can get there in one tap, without knowing or caring which specific event it came from.

Recent shows section: Four event cards showing the most recent shows with upload counts and thumbnail previews. This is a "go deeper" prompt, not the primary content.

Full archive: Collapsed by default behind an "All N shows →" link. When expanded, shows a chronological list of every event grouped by year. This is the archive — real, permanent, fully accessible. It's just not where most visitors start.

The event page in detail

The event page serves two audiences simultaneously: the attendee who was there and wants to find/upload content, and the non-attendee who wants to experience the show.

Watch tab (default): Opens with the most-viewed video for this event playing full-width via Mux player. Below the player: a horizontal scroll of other video clips. If setlist data exists, a "Browse by song" section lists all songs in performance order — each shows a clip count and is tappable to filter the media grid below. This lets a non-attendee step through the show song by song like chapters in a documentary.

Browse tab: Full unfiltered media grid with filter controls: All / Photos / Videos, and section pills (Floor/Pit, Section 100s, Section 200s, Upper deck, Stage left, Stage right). This is for people who want to explore rather than be guided.

Upload tab: The inline upload flow, pre-populated for this event.

The upload flow

The upload experience is the most important UX moment on the platform. If it's too slow or

requires too many steps, fans will post to Instagram instead and Showside gets nothing.

Design constraints:

- From opening the page to upload complete: under 60 seconds on a good mobile connection
- No forced account creation before uploading — ever
- Works in a mobile browser without an app download
- Upload runs in the background — the user shouldn't need to babysit it

The flow:

1. Event confirmation — show entity name, date, venue, city. If arriving from a QR code or direct link, this is pre-confirmed with one tap.
2. File selection — native camera roll picker (photos and videos), drag-and-drop on desktop. Show client-side previews immediately. Limits: photos 20MB each, videos 500MB each, max 10 files per session.
3. Optional tagging — for each file: song (dropdown from setlist), section (pill selector), caption (140 chars). A "Skip all tags" button bypasses this entirely. Tagging is never a gate.
4. Upload — direct to Mux (videos) or Cloudflare R2 (photos) via presigned URLs. Per-file progress bars. "Leave this tab open" message.
5. Success screen — link to the event page, "Share this page" button. Optional account creation prompt: "Want credit for your uploads?" — email + username fields. This is the only place account creation is ever mentioned.

Competitive position

What's been tried:

- CrowdAlbum (2013-2016): Aggregated social posts by location/time data. Acquired by Spotify, shut down. Never tested native direct upload.
- Concert Archives (2013-present): Closest concept. Personal diary for attendees to log concerts. Gets press mentions but sparse media, dated design, no virtual attendance experience for non-attendees.
- FanVids / Aftershow.io / SureShot / Happaning: All B2B/artist-controlled. Fans upload to an artist's marketing CRM. The artist owns the experience. Not a community archive.
- YouTube / TikTok: Fan clips exist and get millions of views but are completely unstructured — no event context, no setlist order, no seat section filtering.

The gap Showside fills: Consumer-first, community-owned, permanently archived, structured for non-attendees. Nobody owns this territory.

The market

- 151 million people attended Live Nation events in 2024
- 56% of concert attendees share media on social during shows (Nielsen Music 360)
- Average top-100 tour ticket price up 50% since 2019 — the priced-out audience grows

every year
- Fan-shot concert clips regularly reach 10M+ views on YouTube — demand is proven at scale
- The "virtual attendance session" — a non-attendee spending 30+ minutes browsing a full show — is the one unproven hypothesis the V1 must test

Tech stack

- **Framework:** Next.js 14 with App Router (SSR required for SEO)
- **Database:** Supabase (Postgres + auth + RLS)
- **Video:** Mux (direct upload, HLS adaptive streaming, webhook processing, thumbnails)
- **Photos:** Cloudflare R2 (presigned PUT URLs, public CDN delivery, zero egress fees)
- **Deployment:** Vercel
- **Styling:** Tailwind CSS
- **Language:** TypeScript throughout

Database overview

Six core tables: `entities`, `events`, `media`, `users`, `likes`, `follows`. All counters (upload_count, like_count, follower_count, photo_count, video_count) are maintained by database triggers in the same transaction as inserts/deletes — not application code. Anonymous users get a session token stored in localStorage; likes and uploads attach to session tokens until/unless the user creates an account. Media moves through a status lifecycle: `uploading` (created on presign) → `active` (photos: on /upload/complete; videos: on Mux webhook firing after processing).

What's been built (V1 status)

- ✅ Full database schema, migrations, RLS policies, triggers, indexes
- ✅ Seed data: 4 entities (Morgan Wallen, Taylor Swift, Philadelphia Phillies, Kentucky Derby), 9 events
- ✅ All lib clients: Supabase (server/client/admin), Mux, R2, session management
- ✅ All API routes: photo upload presign, video upload presign, upload complete, Mux webhook, like toggle, search, QR code generation, admin login
- ✅ Admin panel: dashboard stats, entity CRUD, event CRUD with setlist paste, moderation queue, QR download
- ✅ All public pages: home (discovery feed), entity page (experience-first layout), event page (Watch/Browse/Upload tabs), full-screen media viewer, search, sitemap.xml, robots.txt, JSON-LD structured data
- ⏳ Upload flow UI (Phase 5) — the file picker stepper, direct upload to Mux/R2, success screen

- ⏳ SEO polish, view-count deduplication, infinite scroll, Lighthouse audit (Phase 6)
- ⏳ Design pass — Claude Code built functional UI with its own judgment. A full visual redesign pass is planned once functionality is complete.

What is NOT in V1

- Audio fingerprinting / auto song tagging (V2)
- Mobile app — React Native (V2)
- Automated event ingestion from Ticketmaster / setlist.fm APIs (V2)
- Venue seat map (V2)
- Email notifications (V2)
- Content licensing marketplace (V3)
- Official artist/team partner dashboard (V3)
- Monetization / advertising (V3)

Design direction (for the upcoming design pass)

The visual aesthetic should feel cinematic and fan-authentic — dark backgrounds, warm accent colors, content-forward. It should feel like being in the venue, not like a database browser. Key principles:

- Video is the hero content type — lead with it everywhere
- Dark theme throughout — both because it suits the concert/sports context and because it makes media thumbnails pop
- The entity page hero should feel like a movie poster for the artist — immersive, not corporate
- Filter pills and navigation should feel lightweight and fast — this is a browsing experience, not a research tool
- Mobile experience is primary — most users will arrive on a phone, either uploading after a show or browsing from their couch

The mockup we designed together (visible in prior conversation) shows: dark cinematic hero with video thumbnails on the right, fan highlights grid with mood filters, most filmed songs with relative bars, recent shows cards, and a collapsed archive section at the bottom.

You said: did we set an admin key for vercel?

did we set an admin key for vercel? if so, what is it

May 26

Claude responded: The admin key was set locally in your .

The admin key was set locally in your `.env.local` file — it's whatever you put in for `ADMIN_KEY` when you created that file earlier tonight.

Check it by running this in Terminal:

bash

```
grep ADMIN_KEY ~/Projects/showside/.env.local
```

That'll show you the value you set. You haven't deployed to Vercel yet so it hasn't been added there — that happens when you do the first Vercel deployment and add all your env vars to the Vercel dashboard.